

FDE-ASSOCIATE FOUNDATION CERTIFICATION EXAMINATION

Curriculum Focus: AI Prompting in 2026 (13 Concepts & Operational Harnesses)

Series: 2026-FDE-PR13
35 Continuous Items • 78% Pass
SLA

SINGLE CHOICE Question 1

A client's business team claims that modern AI models operate like an interactive database that stores private conversation history within its internal parameters between turns. During a customer requirements workshop, how should an FDE reframe the fundamental stateless reality of modern LLMs to justify an external context management architecture?

- A. Explain that LLMs preserve memory strictly in dynamic weights that must be re-quantized after each customer session to maintain isolation.
- B. Clarify that the model keeps zero memory of its own between turns and evaluates each request strictly using the text physically present in its context window at that instant; apparent memory is product scaffolding re-injecting stored profiles or transcripts.
- C. Demonstrate that statelessness only applies to open-weights models, whereas proprietary frontier models maintain persistent vector memory tables inside their attention heads.
- D. Recommend increasing sampling temperature to 1.0 so the internal memory registers do not lock into recurrent state loops.

Answer: B

SINGLE CHOICE Question 2

A company executive tests a newly deployed internal marketing assistant by asking: "Write a thought leadership blog post about enterprise microservices." The resulting output contains generic clichés (e.g., 'in today's fast-paced digital landscape', 'a paradigm shift') and lacks actionable substance. Applying Concept 1 (Novice vs Power User), what diagnosis and corrective pattern should the FDE enforce?

- A. The model has suffered catastrophic parameter degradation; switch immediately to an open-weights 8B model.
- B. The prompt is a novice search query producing default AI slop; convert the interaction into a structured briefing loop: provide source docs, outline first, critique the outline, expand into bullet points, and generate prose only after approval.
- C. The temperature was set too low; raising it to 1.6 will remove generic filler text automatically.
- D. The model failed to trigger deep research; prefix the prompt with 'search the entire web thoroughly'.

Answer: B

SINGLE CHOICE Question 3

An FDE is diagnosing an enterprise legal research assistant. When asked about standard international maritime salvage laws, the model provides exhaustive, highly accurate citations. However, when asked about a proprietary internal company joint venture clause written in 2024, it fabricates realistic-looking clause numbers and indemnities. How does Concept 2 (Pretrained Knowledge & The Thickness of the Internet) explain this discrepancy?

- A. Pretrained knowledge is strong where the public internet is thick and dense, but weak or blind where data is private or thin; confidence is statistical probability, not proof of truth.
- B. The model was trained exclusively on pre-2020 Wikipedia archives, disabling modern legal reasoning.
- C. Maritime law relies on code execution, while joint ventures require diffusion generative heads.
- D. The customer prompt triggered sycophancy filters that intentionally corrupt private contract terms.

Answer: A

SINGLE CHOICE Question 4

A financial analytics team wants to understand why their automated analyst sometimes answers within 800ms with general definitions, while other times it pauses for 8 seconds, reads 12 web pages, and other times runs for 3 minutes synthesizing 50 sources. Under Concept 3, what three answering modes operate in modern AI, and what determines which mode fires?

- A. Quantized mode, FP16 mode, and FP32 mode; chosen strictly by available GPU VRAM.
- B. Pretrained knowledge (instant memory recall), web search (looking up fresh facts), and deep research (multi-step investigative exploration across dozens of sources); steered deliberately by prompt wording, task scope, and explicit mode toggles.
- C. Diffusion mode, autoregressive mode, and transformer mode; determined by whether images are uploaded.
- D. Prompt caching mode, KV cache mode, and memory profile mode; selected by the client browser version.

Answer: B

SINGLE CHOICE Question 5

An FDE is reviewing an architectural defect in a long-running customer support agent. By message 45 of an ongoing customer troubleshooting session, the agent begins hallucinating shipping policies that contradict the system instructions established at the start of the chat. According to Concept 4 (Context Window & Context Rot), what is the root cause?

- A. The context window accumulated clutter, resolved issues, and contradictory conversational drift, diluting the model's attention budget across early instructions ('context rot').
- B. The inference server ran out of physical SSD swap space, corrupting token weights.
- C. Extended thinking models automatically flush system prompts after turn 20 to protect user privacy.
- D. The temperature parameter increased proportionally with every turn of conversation.

Answer: A

SINGLE CHOICE Question 6

A software engineering client wants an agent to refactor a complex distributed concurrency lock that involves multi-threaded deadlock edge cases. When prompted in standard mode, the model suggests a naive, buggy mutex approach. What modern 2026 capability (Concept 5) should the FDE activate, and what is its operational mechanism?

- A. LoRA fine-tuning; it permanently patches the base model weights with lock-free algorithms.
- B. Extended thinking (Reasoning mode); the model writes a private chain of internal reasoning tokens to explore multiple candidate strategies, critique dead ends, and verify constraints before generating visible output.
- C. Prompt injection bypass; it disables all safety boundaries to unlock unrestricted execution speeds.
- D. Artifact rendering; it compiles the code into an interactive WebAssembly container in real time.

Answer: B

SINGLE CHOICE Question 7

An FDE is reviewing a client prompt: *"I designed this new enterprise SaaS onboarding workflow. Isn't this one of the most frictionless user experiences you've seen? Please critique it."* The model responds with lavish praise and mild, cosmetic suggestions. According to Concept 6 (Sycophancy), how should the FDE redesign this evaluation prompt?

- A. Increase the temperature to 1.5 to make the model less friendly.
- B. Frame the request neutrally without revealing personal ownership, and provide a strict, scored rubric evaluating failure rates, churn friction, and competitive differentiation.
- C. Add the phrase 'be brutally honest' at the end of the existing prompt without altering the leading question.
- D. Convert the prompt into a diffusion generation script.

Answer: B

SINGLE CHOICE Question 8

In an enterprise team adopting AI for architectural design, engineers frequently generate complete 2,000-word system design proposals in a single monolithic prompt, only to discover major structural flaws in the data layer. According to Concept 7, which iterative operational habit produces higher quality outputs?

- A. The single-shot mega-prompt containing 50 required constraints in one block.
- B. The brainstorm-iterate loop: provide rich context, demand 3 to 5 distinct high-level options, provide targeted human feedback to select a direction, refine options, and only then expand the selected winner.
- C. Running the prompt five times in parallel at temperature 1.0 and taking the longest response.
- D. Converting all design documents into Base64 strings before asking for generation.

Answer: B

SINGLE CHOICE Question 9

A manufacturing client needs an AI solution to analyze warehouse floor hazard photos, transcribe factory supervisor voice memos, and produce structured maintenance logs. Applying Concept 8 (Multimodal), how should the FDE architect the ingestion pipeline?

- A. Deploy three separate single-purpose text-only LLMs connected by shell scripts.
- B. Leverage native multimodal capabilities by passing images, audio memos, and structured text schemas into a unified prompt context window, treating visual and acoustic inputs as native tokens.
- C. Convert images into ASCII art diagrams and audio into Morse code text strings.
- D. Insist that multimodal models cannot process voice memos without external human transcription.

Answer: B

SINGLE CHOICE Question 9

A client product manager wants a quick internal prototype for a regional sales tax calculator. Rather than building a separate full-stack web application with React and Node, what modern chat product feature (Concept 9: Artifacts) should the FDE demonstrate to deliver an interactive tool in one turn?

- A. Prompting the model to generate a self-contained interactive HTML/JS application inside a dedicated side-panel artifact that the PM can test, edit, and download directly.
- B. Writing a Python script that requires Docker compilation on the client's laptop.
- C. Generating raw text pseudocode and instructing the PM to paste it into a spreadsheet formula.
- D. Using diffusion models to render a static JPEG screenshot of a calculator.

Answer: A

SINGLE CHOICE Question 11

An FDE is reviewing an automated financial accounting pipeline. When asked to compute compound interest across 1,500 irregular ledger transactions, a model outputted an approximate value of \$142,300, which differed from the true audit amount of \$138,412.19. Applying Concept 10 (Code Execution), what was the failure and what is the fix?

- A. The model hallucinated due to low temperature; increase temperature to 0.9.
- B. The model estimated the calculation via token generation rather than running code; the prompt must explicitly command: 'Write a Python program, execute it on this dataset, and report the calculated output.'
- C. The CSV dataset was too clean, preventing the attention heads from calculating variance.
- D. Code execution is an obsolete 2022 feature that has been replaced by diffusion math.

Answer: B

SINGLE CHOICE Question 12

An enterprise deploys autonomous AI desktop applications (e.g., Cowork, OpenWork) across corporate laptops. A user grants the desktop agent full read/write/execute permissions on their entire root filesystem. According to Concept 11 (AI Desktop Apps and Safety), what core security principle has been violated?

- A. The principle of token conservation.
- B. The principle of least privilege: AI desktop agents should only be granted scoped access to specific project working folders and require explicit confirmation for destructive operations.
- C. The principle of multimodal parity.
- D. The requirement that desktop agents only run inside web browser iframes.

Answer: B

SINGLE CHOICE Question 13

A client insists on standardizing their entire enterprise AI infrastructure on a single provider's flagship model for every task: routing, complex legal analysis, simple data parsing, and image tagging. Why does Concept 12 (The Model Ladder & Jagged Frontier) warn against this monolithic approach?

- A. Provider APIs intentionally rate-limit single-model architectures after 1,000 requests.
- B. Capabilities, speeds, and per-token costs vary widely across model tiers; optimal architectures use a model ladder, routing simple tasks to fast/inexpensive models and reserving heavy reasoning models for multi-step complex work.
- C. Using a flagship model for simple parsing invalidates enterprise compliance certificates.
- D. Frontier models cannot parse JSON or plain text.

Answer: B

SINGLE CHOICE Question 14

An enterprise needs to evaluate which of two candidate LLMs produces superior executive summaries of complex biomedical research papers, but no human medical doctors are available to score thousands of daily outputs. According to Concept 13 (Model-Graded Evals & Cross-Family Grading), what is the most reliable automated evaluation protocol?

- A. Have Model A grade its own output on a 1–5 scale to verify its confidence.
- B. Use a smaller model from the same model family as Model A to grade Model A.
- C. Implement cross-family model evaluation: have independent frontier models from different companies grade the outputs against a strict, objective factual rubric and compare scores.
- D. Count the total number of characters and select the longest summary.

Answer: C

SINGLE CHOICE Question 15

During customer discovery, a client's legal counsel raises concerns that customer support conversations might be used to retrain the provider's public models. What architectural configuration must an FDE mandate to satisfy enterprise confidentiality?

- A. Add an instruction in the prompt: 'Do not remember this conversation for training.'
- B. Migrate the workload to commercial enterprise API endpoints with zero data retention (ZDR) agreements and explicit terms barring training on customer payloads.
- C. Switch the temperature parameter to 0.0 to prevent data logging.
- D. Encode all customer prompts in binary hexadecimal before sending.

Answer: B

SINGLE CHOICE Question 16

An FDE is testing a command-line agent (such as Claude Code or OpenCode) configured inside a customer's software repository. The agent makes repeated syntax errors across five files because it lacks awareness of the project's custom TypeScript types. What context injection technique fixes this at session launch?

- A. Running the agent with the `--force` flag.
- B. Providing a concise project context brief (``CLAUDE.md`` / ``AGENTS.md``) at repository root containing build commands, architecture guidelines, and key typing conventions.
- C. Deleting the Git history to reduce repository token weight.
- D. Disabling code execution in the CLI harness.

Answer: B

SINGLE CHOICE Question 17

In a customer workflow, users repeatedly ask the assistant to summarize 100-page policy PDFs. The system often misses crucial clauses located between pages 40 and 60. How should an FDE configure the document ingestion pipeline to counter positional attention degradation (lost-in-the-middle)?

- A. Chunk the document into logical sections, index them via targeted retrieval, and place key instructions and relevant excerpts at the end and beginning of the prompt context.
- B. Invert the PDF upside down before OCR processing.
- C. Double the sampling temperature to force attention to scan middle tokens.
- D. Remove all punctuation from the policy document.

Answer: A

SINGLE CHOICE Question 18

An autonomous research assistant is given a broad mandate: *"Research the latest semiconductor export regulations."* It enters a cycle where it searches Google, reads a blog post, searches again for a term mentioned in the post, and continues searching indefinitely without generating a brief. What prompt harness constraint is missing?

- A. A diffusion model pipeline.
- B. An explicit termination criterion, fixed step budget (e.g., max 5 tool calls), and a structured output schema requiring final synthesis.
- C. A 1,000,000 token context window upgrade.
- D. Setting temperature to 0.0 to stop all tool calling.

Answer: B

SINGLE CHOICE Question 19

A client asks an FDE why the AI model sometimes refuses benign questions containing security terminology (e.g., 'How does a SQL injection vulnerability work in an educational sandbox?'). Where does this refusal behavior originate, and how should an FDE guide the client?

- A. The refusal comes from the system prompt and RLHF safety training layers; the client should reframe the prompt with explicit defensive educational context and authorization boundaries.
- B. The model's database crashed due to SQL syntax in the user query.
- C. Refusals are caused by corrupted byte-pair tokenizers.

MULTI-SELECT (CHOOSE 2) Question 20

An FDE is conducting an architectural review of an agentic workflow that frequently generates plausible but completely unverified factual statements when drafting corporate press releases. Which TWO interventions directly eliminate this failure mode?

- ☐ Forcing the model to ground its statements exclusively in verified reference documents passed into the context window, with an explicit instruction to state 'Source not provided' for missing facts.
- ☐ B. Increasing the temperature to 1.4 to expand creative vocabulary.
- ☐ C. Enabling live web search or deep research tools to fetch and cite external corroborating URLs before drafting.
- ☐ D. Disabling code execution in the agent harness.
- ☐ E. Using a smaller 1B model to minimize generative hallucinations.

Correct answers: A and C

MULTI-SELECT (CHOOSE 2) Question 21

A client enterprise notices that their AI support workspace experiences severe latency spikes and declining answer quality as employees keep adding questions to single, month-long chat sessions. Which TWO engineering practices counteract this context degradation?

- ☐ A. Encouraging users to start fresh chats for new, unrelated tasks rather than maintaining monolithic threads ('cleaning the desk').
- ☐ B. Implementing automatic context compacting or rolling summarization of older dialogue turns in the API orchestrator.
- ☐ C. Increasing the top-p sampling parameter to 1.0.
- ☐ D. Converting all incoming user text into audio prompts.
- ☐ E. Permanently disabling prompt caching across the API cluster.

Correct answers: A and B

MULTI-SELECT (CHOOSE 2) Question 22

When evaluating a proposed corporate AI strategy, an FDE identifies that the team relies exclusively on novice prompting techniques. Which TWO indicators in their prompt repository demonstrate novice behavior rather than power-user briefing habits?

- ☐ A. Prompts consist of single, generic questions without accompanying reference files, constraints, or specified target audiences.
- ☐ B. Prompts provide structured evaluation rubrics and explicitly ask for 3 to 5 distinct options before final expansion.
- ☐ C. Prompts ask the model leading questions seeking approval (e.g., 'Do you think my marketing plan is good?'), inviting sycophantic flattery.
- ☐ D. Prompts require the model to execute a Python script to verify numerical outputs.
- ☐ E. Prompts use dedicated project workspaces with standing instructions and domain documentation.

Correct answers: A and C

MULTI-SELECT (CHOOSE 2) Question 23

An FDE is designing an AI development platform for financial analysts processing sensitive corporate earnings spreadsheets. Which TWO capabilities are essential to ensure numerical precision and verifiable audit trails?

- ☐ Built-in code execution where the model writes and runs Python scripts to compute exact sums and ratios rather than estimating.
- ☐ B. Side-panel interactive artifacts that allow analysts to inspect formulas, review generated tables, and export clean CSVs.
- ☐ C. Setting temperature to 1.8 to ensure creative statistical modeling.
- ☐ D. Relying solely on pretrained model weights without uploading source spreadsheets.
- ☐ E. Generating raw audio files instead of structured tabular data.

Correct answers: A and B

MULTI-SELECT (CHOOSE 2) Question 24

An FDE is configuring an enterprise Project workspace (such as Claude Projects or custom workspace harnesses) for a team of 50 technical underwriters. Which TWO items represent the ideal standing assets to load into the workspace context layer?

- ☐ A. Standardized enterprise underwriting guidelines, policy rate cards, and approved template schemas.
- ☐ B. Standing instructions defining persona tone, mandatory disclaimers, and required output formatting.
- ☐ C. Raw video recordings of all internal company all-hands meetings from the past 5 years.
- ☐ D. Unencrypted production database passwords and master API keys.
- ☐ E. A complete copy of the Linux kernel source code.

Correct answers: A and B

MULTI-SELECT (CHOOSE 2) Question 25

A client asks an FDE how to verify the factual claims of an AI-generated competitive landscape report before presenting it to the board of directors. Which TWO verification actions should the FDE instruct the client to perform?

- ☐ A. Spot-check external citations and URLs by opening them directly to ensure the cited sources actually contain the quoted facts.
- ☐ B. Submit the draft report to an independent frontier model from a different model family with a strict factual verification prompt.
- ☐ C. Ask the original model: 'Did you make any mistakes?' in the same chat thread.
- ☐ D. Assume that because the model generated footnotes, all facts are automatically legally verified.
- ☐ E. Re-run the prompt at temperature 0.0 and assume identical wording proves factual truth.

Correct answers: A and B

MULTI-SELECT (CHOOSE 2) Question 26

An FDE is conducting security threat modeling on desktop AI agents that possess local filesystem access. Which TWO attack vectors or operational risks must be mitigated?

- ☐ A. Indirect prompt injection where a malicious downloaded text or HTML file instructs the agent to exfiltrate private credentials.
- ☐ B. Unintentional deletion or overwriting of sensitive files when the agent is granted unrestricted directory write permissions.
- ☐ C. The agent exceeding token quotas by reading plain text configuration files.
- ☐ D. The agent failing to render diffusion images on CPUs.
- ☐ E. The agent converting CSV data into markdown tables.

Correct answers: A and B

MULTI-SELECT (CHOOSE 2) Question 27

An FDE is training engineers to build production prompts that reliably invoke reasoning modes when tackling difficult algorithmic tasks. Which TWO prompt structures reliably trigger deep, multi-step problem solving?

- ☐ A. Explicitly prompting the model to explore multiple candidate strategies, evaluate potential edge cases, and verify constraints before writing the final code.
- ☐ B. Selecting an extended thinking / reasoning model mode and setting a generous thinking token budget.
- ☐ C. Adding 'ASAP' and 'urgent' to the prompt text.
- ☐ D. Constraining the output length to exactly 50 words.
- ☐ E. Using only capital letters throughout the prompt instructions.

Correct answers: A and B

TABLE ANALYSIS Question 28

An FDE is evaluating four prompt engineering approaches for an automated customer dispute triage assistant. The business demands: (1) sycophancy score under 10%, (2) factual accuracy over 95%, (3) turnaround latency under 2.5s, and (4) structured JSON output. Review the experimental trial results below:

Method	Prompt Structure & Strategy	Sycophancy Score	Factual Accuracy	Latency	Output Format
M1	Open-ended question ("Do you agree this claim is valid?")	78%	64%	1.1s	Free-form prose
M2	Neutral framing + 5-point rubric + constrained JSON schema	4%	97%	1.8s	Validated JSON
M3	Deep Research mode + 10 web searches	6%	96%	48.0s	Markdown report
M4	Few-shot examples with leading confirmation bias	62%	71%	1.6s	Validated JSON

Which method strictly satisfies all four business requirements?

- A. Method M1
- B. Method M2
- C. Method M3
- D. Method M4

Answer: B

TABLE ANALYSIS Question 29

An FDE audits token consumption and response latency across four enterprise chat configurations running within a customer's corporate workspace over a 30-turn user session:

Session ID	Context Management Strategy	Turn 30 Context Size	Turn 30 Latency	Context Rot Incidents
S-101	Unmanaged monolithic thread (no compaction)	184,000 tokens	14.2s	8 policy violations logged
S-102	Rolling summarization (last 5 turns + summary)	12,400 tokens	1.9s	0 policy violations logged
S-103	Fresh chat per subtask + standing Project instructions	8,500 tokens	1.4s	0 policy violations logged
S-104	Unmanaged thread + prompt caching enabled	184,000 tokens	4.6s	7 policy violations logged

What critical operational conclusion should the FDE report to the client regarding session S-104?

- A. Prompt caching eliminated all context rot, proving that large context windows require no management.
- B. Prompt caching reduced latency but failed to prevent context rot; structural context management (compaction or fresh chats) is mandatory to maintain instruction fidelity.
- C. Session S-104 crashed because prompt caching is incompatible with 30-turn dialogues.
- D. Rolling summarization consumes more tokens than unmanaged monolithic threads.

Answer: B

TABLE ANALYSIS **Question 30**

An FDE evaluates four model configurations on a complex financial reconciliation task requiring arithmetic verification across 500 bank statements:

Config	Model Tier	Tool Access	Arithmetic Accuracy	Execution Time
CFG-A	Fast / Small LLM (8B)	None (Direct token generation)	31.4%	0.8s
CFG-B	Frontier Reasoning LLM	None (Extended thinking CoT)	79.2%	12.4s
CFG-C	Fast / Small LLM (8B)	Code Execution (Python interpreter)	100.0%	1.6s
CFG-D	Frontier LLM	Web Search	42.0%	8.1s

What fundamental FDE implementation principle does this telemetry demonstrate regarding arithmetic and data manipulation?

- A. Reasoning models should always be chosen over code execution for financial reconciliation.
- B. For deterministic computation on structured data, a fast model equipped with code execution outperforms even frontier reasoning models in accuracy, latency, and cost.
- C. Web search is necessary to verify internal ledger balances.
- D. Code execution produces frequent floating-point errors compared to text generation.

Answer: B

TABLE ANALYSIS **Question 31**

An FDE conducts model-graded evaluations for an automated compliance summarizer. Outputs from Model X and Model Y were evaluated by three different evaluator configurations:

Evaluation Run	Evaluator Model	Relationship to Target	Model X Score	Model Y Score	Observed Evaluator Bias
Eval-1	Model X	Self-evaluation	94/100	71/100	Severe self-preference bias
Eval-2	Model X-Mini	Same model family	91/100	73/100	Family-level stylistic bias
Eval-3	Model Z (Competitor)	Cross-family independent	82/100	84/100	Objective rubric alignment
Eval-4	Human Blind Panel	Ground-truth gold standard	81/100	85/100	No model bias

Why is Eval-3 considered the industry-standard automated grading pattern (Concept 13) when human experts are unavailable?

- A. Same-family models share blind spots and exhibit mutual self-preference; cross-family grading breaks familial bias and correlates closely with human gold standards.
- B. Model Z is always cheaper to run than Model X.
- C. Competitor models are programmed to give higher scores to other companies' outputs.
- D. Self-evaluation is mathematically proven to be 100% accurate.

Answer: A

TABLE ANALYSIS **Question 32**

An FDE profiles the impact of four different prompt briefing formats on an AI code refactoring workflow tasked with restructuring legacy Python modules:

Briefing ID	Provided Briefing Artifacts	Task Specification	Syntax Pass Rate	Regression Test Failures
BR-1	None ("Make this Python code cleaner")	Vague single ask	62%	14 regressions
BR-2	Target file only	"Refactor into classes"	79%	8 regressions
BR-3	Target file + test suite + type definitions	"Refactor with zero test breakage"	98%	0 regressions
BR-4	Target file + 10 unrelated project files	"Refactor this file"	84%	5 regressions

Which prompt briefing strategy yielded zero regressions, and what principle does it illustrate?

- A. BR-1; vague instructions give the model complete creative freedom.
- B. BR-3; power-user prompting provides relevant context (tests, types) and clear constraints, allowing the model to self-verify against ground truth.
- C. BR-4; overloading unrelated context improves code refactoring.
- D. BR-2; classes are always superior to functions in AI generation.

Answer: B

TABLE ANALYSIS **Question 33**

An FDE tests four prompt framing strategies on a creative strategy team that frequently complains about getting generic "yes-man" praise from AI assistants:

Prompt Framing	Prompt Template Text	Critical Flaws Identified	Actionable Improvements	User Feedback
Framing-A	"Here is my product pitch deck. What do you think? I spent months on it."	0 flaws	2 minor compliments	"Useless sycophantic flattery"
Framing-B	"Critique this pitch deck as an adversarial VC evaluating an investment."	4 flaws	6 strategic fixes	"Highly actionable and rigorous"
Framing-C	"Be very honest and tell me if this is good or bad."	1 flaw	3 general tips	"Still leaned toward agreeing"
Framing-D	"Rate this pitch deck 10/10."	0 flaws	0 fixes	"Pure praise"

Why was Framing-B significantly more effective than Framing-C at overcoming sycophancy?

- A. Framing-B assigned an adversarial role with an explicit incentive structure (VC evaluating investment risk), removing conversational politeness bias.
- B. Framing-B automatically increased inference temperature to maximum.
- C. Framing-C crashed the model's tokenizer.
- D. Framing-B invoked deep research web search tools automatically.

Answer: A

TABLE ANALYSIS **Question 34**

An FDE is conducting security audits of desktop AI applications across four employee workstations:

Workstation	Desktop AI App	Permissions Granted	File Scope	Audit Log Event
WS-1	Cowork Desktop	Full Read / Write / Exec	Root (`/`)	Agent accidentally overwrote system hosts file during build
WS-2	OpenWork	Read Only	Dedicated `/project/docs`	Clean execution, 0 incidents
WS-3	Local Agent CLI	Read / Write with Confirmation	Working git repo folder	User blocked accidental deletion of `.env` secret file
WS-4	Custom Agent	Full Read / Write	Entire user home directory	Agent committed unencrypted private SSH keys to repo

What key FDE safety rule from Concept 11 is substantiated by the security incidents on WS-1 and WS-4?

- A. Desktop AI tools should never be used under any circumstances.
- B. Desktop tools must be confined to the smallest directory scope the job needs, with interactive confirmation gates for write/delete actions.
- C. SSH keys should be shared in the system prompt to avoid permission errors.
- D. Home directories must always be formatted in FAT32 to prevent agent access.

Answer: B

TABLE ANALYSIS **Question 35**

An FDE benchmarks four model routing tiers for an enterprise customer handling 100,000 daily customer queries across three distinct complexity tiers (Level 1: FAQs; Level 2: Invoice data extraction; Level 3: Multi-contract legal dispute resolution):

Strategy ID	Level 1 Routing (60k reqs)	Level 2 Routing (30k reqs)	Level 3 Routing (10k reqs)	Daily Cost & SLA Status
STRAT-1	Frontier Reasoning Model	Frontier Reasoning Model	Frontier Reasoning Model	\$14,200 / day (Budget breached by 300%)
STRAT-2	Fast/Lightweight Model	Fast/Lightweight Model	Fast/Lightweight Model	\$420 / day (Failed 62% of Level 3 legal tasks)
STRAT-3	Fast/Lightweight Model	Medium Model + Code Exec	Frontier Reasoning Model + RAG	\$1,850 / day (100% SLA pass rate across all tiers)
STRAT-4	Deep Research Mode	Deep Research Mode	Deep Research Mode	\$35,000 / day (Latency exceeded 60s SLA on all reqs)

Which strategy implements the optimal "Model Ladder" architecture (Concept 12), balancing cost, speed, and capability?

- A. STRAT-1
- B. STRAT-2
- C. STRAT-3
- D. STRAT-4

Answer: C