

FDE-Associate Foundation — AI Prompting 2026

35 Scenario-Based MCQs | 19 single-answer + 8 two-answer + 8 table-based

Source basis: The AI Agent Factory / Panaversity, "AI Prompting in 2026: A Crash Course." Questions are primarily grounded in the source and extended into FDE-style discovery, architecture, troubleshooting, evaluation, reliability, security, and delivery scenarios.

1. A retail customer asks an FDE to build an AI-assisted vendor-selection workflow. The customer says, "Compare these vendors and recommend the best one." The FDE has three vendor quotes, an internal cost spreadsheet, the procurement policy, and notes that the buyer values implementation speed over lowest price. Which initial prompt is most likely to produce an actionable result?

- A. "Which vendor is best? Give me your recommendation."
- B. "Compare the vendors. Be detailed and professional."
- C. "Using the attached quotes, cost sheet, and procurement policy, recommend a vendor. Goal: minimize total business risk while prioritizing implementation speed. Limits: comply with procurement policy and identify unsupported assumptions. Ask: produce a comparison, key trade-offs, and the conditions that would flip the recommendation."
- D. "Read all files and tell me everything important about the vendors."

Answer: C

2. A customer asks whether a newly released regulation permits a proposed AI workflow. The team has a strong internal summary written six months ago, but the regulation may have changed last month. Which approach best matches the source's retrieval guidance?

- A. Use pretrained knowledge because the internal summary is already detailed.
- B. Use the internal summary only and avoid web search to reduce noise.
- C. Use web/search against named authoritative sources, require citations for claims, and mark unsupported claims as unverified.
- D. Ask the model to reason harder from the old summary until it reaches high confidence.

Answer: C

3. An FDE is troubleshooting an AI workflow that gives increasingly generic answers after a 70-message conversation covering requirements, coding, travel planning, and unrelated writing tasks. The model still has a large context window. What is the best first intervention?

- A. Add every previous file to the current prompt so the model has more context.
- B. Start a new conversation scoped to the current task and provide only relevant context.
- C. Ask the model to summarize the entire 70-message history before answering.
- D. Turn on extended thinking and keep the same conversation.

Answer: B

4. A customer says, "We uploaded the contract last week, so the model should remember the exact termination clause today." The product has memory enabled. Which explanation is most accurate?

- A. Memory means the model permanently stores all previous chat text.
- B. The model itself becomes stateful once memory is enabled.
- C. Memory is a product-level note placed into the model's current context; the model remains stateless, so exact prior document details should be supplied or retrieved again.
- D. Memory guarantees that every uploaded file remains available in every new conversation.

Answer: C

5. An FDE is asking an AI to analyze 12,000 support records. The prompt contains the business goal, schema, metric definitions, and 18 unrelated historical exports. The answer becomes less precise. What is the best diagnosis?

- A. The model needs more unrelated examples to generalize.
- B. The problem is necessarily model intelligence, not prompt design.
- C. Context curation is poor: duplicate/off-topic material is increasing the chance of misreading relevant evidence.
- D. The solution is to remove all metric definitions and let the model infer them.

Answer: C

6. A product manager asks an AI for five campaign ideas but specifies the exact structure, slogan, tone, word count, channel mix, and sentence pattern for every idea. The five outputs are nearly identical. What should the FDE change first?

- A. Increase the requested number from five to fifty while keeping every constraint.
- B. Remove true business constraints such as audience and prohibitions.
- C. Loosen output-shape constraints during brainstorming, then restore precision after selecting a direction.
- D. Ask for a longer reasoning trace before generating ideas.

Answer: C

7. A customer wants an objective critique of an internal proposal. Their prompt says, "This is clearly the right approach; explain why it will work." Which redesign most directly reduces sycophancy?

- A. Ask the model to be more enthusiastic and persuasive.
- B. Ask, "Is this good?" and request a confidence score.
- C. Reframe neutrally and provide named criteria scored independently, with one improvement per criterion.
- D. Provide more background text while keeping the conclusion in the prompt.

Answer: C

8. An FDE is deciding whether to use a brainstorming loop or directly request a final implementation specification. Requirements are still ambiguous and the customer has not selected among competing solution directions. Which sequence is strongest?

- A. Draft the final specification immediately so the customer can react to something concrete.
- B. Generate materially different options, collect customer feedback, iterate, then tighten constraints around the chosen direction.
- C. Ask for one option and repeatedly improve it until it reaches 9.5.
- D. Remove all context until the model produces a more creative answer.

Answer: B

9. A difficult architecture decision has two plausible options with different latency, cost, and reliability implications. The FDE wants the model to expose the trade-offs rather than give a superficial answer. Which prompt is best aligned with the source?

- A. "Pick the best architecture."
- B. "Think hard. Identify the three trade-offs that matter, recommend one, and state the conditions under which your recommendation would flip."
- C. "Explain both architectures in maximum detail without choosing."
- D. "Assume Option A is better and justify it rigorously."

Answer: B

10. A team asks an AI to calculate the median, average, and outliers in a dataset. The model returns plausible numbers but shows no evidence that code executed. What is the most important next step?

- A. Accept the result because arithmetic is simple.
- B. Ask the model to rewrite the explanation more confidently.
- C. Explicitly require it to write and run code and show the code/result used for the calculation.
- D. Replace the dataset with a larger one.

Answer: C

11. A customer wants an AI-generated clinical note from a 45-minute recorded consultation. The audio contains technical terms and several uncertain statements. Which workflow is safest and most useful?

- A. Generate a polished note without uncertainty markers so it reads professionally.
- B. Transcribe and structure the note, explicitly flagging low-confidence portions, while leaving clinical judgment to the professional.
- C. Use live voice conversation because it is more natural.
- D. Summarize the audio from memory after listening once, without checking unclear sections.

Answer: B

12. An FDE is building a one-screen internal calculator with no accounts and no external services. The customer can describe the desired behavior clearly. Which implementation approach is most consistent with the source?

- A. Treat it as a large software engineering project requiring a full multi-service architecture before prototyping.
- B. Use a Goal/Input/Output brief to generate a small artifact, then test and iterate it.
- C. Ask for a generic website template and manually infer the missing behavior.
- D. Use an agent to redesign the entire customer platform first.

Answer: B

13. A generated app works in the chat artifact but must support internet multiplayer, accounts, and real-time matchmaking. What should the FDE conclude?

- A. One-prompt artifacts are designed for this exact scale if the prompt is long enough.
- B. The artifact should be repeatedly regenerated until networking emerges.
- C. The request has crossed beyond the source's small one-screen/no-external-services sweet spot and needs real engineering beyond a one-prompt build.
- D. The correct fix is to remove the Goal/Input/Output structure.

Answer: C

14. A customer wants a generated dashboard to be "better." The FDE has only a vague complaint and no acceptance criteria. Which next step most improves the feedback loop?

- A. Ask the model to decide what "better" means.
- B. Provide specific observed failures and desired behavior, then score the dashboard against named criteria.
- C. Add more unrelated customer documents.
- D. Ask the model to regenerate the dashboard from scratch without feedback.

Answer: B

15. An FDE has a high-stakes proposal. The primary model scores its own draft 9.5/10. No independent model or human has reviewed it. Which conclusion is most defensible?

- A. The 9.5 proves the proposal is correct.
- B. The score is useful as a progress signal, but not equivalent to independent verification.
- C. A higher self-score can replace source checking.
- D. The proposal should be shipped because self-critique eliminates model blind spots.

Answer: B

16. A confidential customer document is covered by an NDA. The FDE wants to use cross-model review with several consumer AI products. What is the best decision criterion before pasting the document into another tool?

- A. Choose the model with the highest public leaderboard score.
- B. Choose the fastest model regardless of data policy.
- C. Check each tool's data/privacy policy and only use services consistent with the confidentiality obligations.
- D. Remove the title page; that alone makes the document safe.

Answer: C

17. A team keeps correcting the same wrong assumption in every new chat. The prompts are otherwise well designed. The product has a persistent memory/profile feature. What should the FDE investigate?

- A. Increase the context window.
- B. Rewrite the same correction in every chat forever.
- C. Inspect and correct the stored memory/profile because the repeated assumption may be coming from that layer.
- D. Disable all reasoning features.

Answer: C

18. A customer asks for an AI research report. They specify that the report must prove their preferred hypothesis and provide supporting sources only. What is the best FDE response pattern?

- A. Keep the requested hypothesis as the conclusion and search only for supporting evidence.

- B. Treat the desired conclusion as an assumption, loosen the conclusion during research, define evidence standards and sources, and let findings determine the result.
- C. Avoid sources because source constraints bias the model.
- D. Ask for three persuasive arguments for the hypothesis before doing any research.

Answer: B

19. A customer's support-analysis workflow is producing confident but weakly grounded recommendations. Which TWO interventions are most directly aligned with the source?

- A. Add every historical ticket and attachment to maximize context volume.
- B. Define the relevant evidence sources and require unsupported claims to be marked as unverified.
- C. State the desired conclusion so the model has a clearer target.
- D. Use a rubric with explicit criteria and score each criterion separately.
- E. Keep all unrelated conversations in the same thread to preserve continuity.
- F. Remove the customer's source material so pretrained knowledge can dominate.

Answer: B and D

20. An FDE is designing an iterative solution-selection workshop. Which TWO practices best implement the source's brainstorm-iterate loop?

- A. Ask for multiple materially different options before selecting one.
- B. Fully specify the final implementation before any options are generated.
- C. Reject selected options with reasons and request new options incorporating the feedback.
- D. Keep the first option and ask for increasingly polished wording.
- E. Avoid customer feedback until the final draft to reduce churn.
- F. Require every brainstorm option to use the exact same structure and length.

Correct answers: A and C

21. A long-running FDE investigation has begun mixing unrelated workstreams. Which TWO actions are most appropriate?

- A. Start a new scoped conversation when the topic changes materially.
- B. Add even more historical messages to restore all details.
- C. Remove irrelevant and duplicate attachments from the active context.
- D. Assume a large context window prevents context rot.
- E. Ask the model to infer missing constraints from older unrelated tasks.
- F. Disable all customer-specific context permanently.

Correct answers: A and C

22. A team wants stronger evaluation of a production-facing draft. Which TWO practices are explicitly supported by the source?

- A. Ask for named criteria scored 1–10 with a justification.
- B. Accept "looks good" as sufficient evaluation.
- C. Have a second model from a different family apply the same rubric.
- D. Use only models from the same family to ensure consistent judgment.
- E. Treat the score as a guarantee of factual correctness.
- F. Avoid independent review when the main model scores above 9.

Correct answers: A and C

23. An FDE is evaluating a spreadsheet-analysis result where the output is numerically plausible. Which TWO actions best address the source's silent code-execution failure mode?

- A. Explicitly ask the model to write and run code for the calculation.
- B. Accept plausible numbers if the prose explanation is confident.
- C. Request/show the code used so execution can be inspected.

- D. Ask for more adjectives describing the trend.
- E. Replace exact calculations with estimates.
- F. Remove the dataset and ask for a general explanation.

Correct answers: A and C

24. A customer wants to deploy a small generated web tool. Which TWO conditions indicate the one-prompt artifact approach is a good fit?

- A. It fits on one screen and has no account system.
- B. It requires internet multiplayer and real-time matchmaking.
- C. It has no external services.
- D. It requires a distributed backend with several authentication providers.
- E. It requires cross-region state synchronization.
- F. It must replace the customer's entire production platform.

Correct answers: A and C

25. A medical customer wants to use mixed inputs for an AI workflow. Which TWO design choices best follow the source?

- A. Combine relevant inputs such as a transcript, document, image, or voice memo when each contributes evidence.
- B. Treat all audio interpretation as fully reliable, including sarcasm and speaker attribution.
- C. Flag uncertain transcription or interpretation instead of silently inventing content.
- D. Prefer live voice for every technical task because it is faster.
- E. Use general AI for non-speech music analysis whenever possible.
- F. Remove human judgment from the workflow once the multimodal model responds.

Correct answers: A and C

26. A high-stakes customer decision is being checked across models. Which TWO statements correctly reflect the source's limits?

- A. Different model families can expose different blind spots.
- B. Three models agreeing proves the answer is true.
- C. Cross-model checking is a progress signal, not a truth guarantee.
- D. Human review is unnecessary once two independent models score the work above 9.
- E. Models from different companies always have completely independent training knowledge.
- F. Cross-model checking eliminates the need to inspect data/privacy policies.

Correct answers: A and C

27. An FDE must choose a prompting strategy for four customer tasks. The customer wants the strategy that best fits the task type rather than one universal prompt style.

['Task', 'Current request', 'Relevant context', 'Desired behavior', 'Observed risk']

['Brainstorm', 'Generate launch directions', 'Audience + product + things to avoid', 'Diverse options first', 'Near-duplicate ideas']

['Research', 'Map evidence on a disputed claim', 'Named sources + evidence standard', 'Let evidence determine conclusion', 'Confirmation bias']

['Drafting', 'Produce final customer memo', 'Voice + facts + format + length', 'One precise deliverable', 'Fluent wrong draft']

['Analysis', 'Explain what sales data says', 'Data + metric definitions + exact question', 'Neutral interpretation', 'Flattering conclusion']

- A. Brainstorm, because output structure should be loosened while true problem constraints remain.
- B. Research, because source requirements should be removed.
- C. Drafting, because all constraints should be removed until after the first draft.
- D. Analysis, because the conclusion should be supplied explicitly.

Correct answers: A

28. A customer reports that a long AI-assisted incident investigation has degraded. The FDE records the following evidence.

Evidence	Observation	Likely mechanism	Best intervention	Why
A	New unrelated topics were added for 40+ turns	Context rot	Start a new scoped conversation	Removes stale unrelated working context
B	The model repeatedly contradicts an old constraint	Compaction may have lost detail	Restate only the active constraints	Restores required detail
C	The same wrong personal assumption appears in new chats	Persistent memory/profile	Inspect and correct stored memory	Fixes repeated cross-chat assumption
D	A relevant file is duplicated four times	Poor context curation	Remove duplicates	Reduces ambiguity/noise
E	The model lacks a current regulation	Knowledge cutoff	Use web/search	Retrieves current evidence

A. A. A

B. B. B

C. C. C

D. D. D

Answer: C

29. An FDE evaluates a generated data-analysis response. The goal is to identify whether the response demonstrates actual computation.

Run	Prompt	Response evidence	Numerical result	Execution evidence
1	Compute median and mean	Confident prose only	Plausible	None
2	Compute median and mean; write/run code	Code block + result	Plausible	Visible code
3	Compute from uploaded CSV	Narrative summary	Plausible	None
4	Compute and show code used	Code + result	Matches dataset	Visible code

A. A. Run 1

B. B. Run 2

C. C. Run 3

D. D. Run 4

Answer: D

30. A customer asks for an objective review of a generated proposal. The FDE runs several evaluation passes.

Pass	Evaluator	Method	Score	Limitation
1	Primary model	"Is this good?"	9.5	High flattery risk
2	Primary model	Named rubric, 1–10 per criterion	9.1	Self-evaluation
3	Different model family	Same rubric	8.6	Independent model, still imperfect
4	Human domain expert	Checks critical factual claims	Accept/modify	Required for high-stakes facts

A. A. Pass 1 is strongest because it has the highest score.

B. B. Pass 2 proves the proposal is correct because it uses a rubric.

C. C. Pass 3 adds independent signal, while Pass 4 is necessary for critical factual validation.

D. D. Pass 4 is unnecessary because cross-model review is more objective.

Answer: C

31. An FDE is selecting a build approach for a customer request.

Request	Scope	External dependencies	Accounts/networking	Best fit
A	Single-screen timer	None	None	One-prompt artifact
B	Simple calculator	None	None	One-prompt artifact
C	Multiplayer game	Real-time backend	Required	Real engineering
D	Internal tool with complex integrations	Several APIs	Required	Beyond one-prompt artifact

- A. A. A
- B. B. C
- C. C. D
- D. D. Both A and B

Answer: D

32. A team is comparing model-review strategies for a sensitive customer project.

Strategy	Models	Data handling check	Review depth	Primary use
A	One model	Not checked	Quick	Low-stakes drafting
B	One model	Checked	Rubric + iteration	Routine quality improvement
C	Two families	Checked	Same rubric	Higher-confidence review
D	Three families	Checked	Cross-model debate	High-stakes review

- A. A. A
- B. B. B
- C. C. C
- D. D. D

Answer: D

33. An FDE is preparing a customer-facing research workflow. The team has four possible evidence strategies.

Strategy	Source selection	Conclusion handling	Citation behavior	Main risk
A	Unnamed blogs	Preselected conclusion	None	Unsupported claims
B	Named authoritative sources	Evidence determines result	Citation per claim	Lower grounding risk
C	Internal memory only	Preselected conclusion	None	Stale/biased evidence
D	General web search	Desired conclusion in prompt	Mixed	Confirmation bias

- A. A. A
- B. B. B
- C. C. C
- D. D. D

Answer: B

34. A customer is building a personal professional site with AI. The first version is generic; the FDE then collects stronger evidence.

Iteration	Inputs provided	AI behavior	Quality issue	Next move
1	"Build a site about me"	Generic sections	Average-person assumptions	Add goal, audience, decisions
2	Goal + audience + design choices	More specific site	Missing proof of skills	Attach certificate and concrete project evidence
3	Certificate + shipped project + course knowledge	Evidence-backed content	Needs prioritization	Refine against visitor clarity

Iteration	Inputs provided	AI behavior	Quality issue	Next move
4	Visitor-oriented feedback	Targeted edits	Remaining polish gaps	Iterate and verify

- A. A. 1
- B. B. 2
- C. C. 3
- D. D. 4

Answer: C

35. A customer's AI workflow must answer a question about a current rule. The FDE compares four prompt designs.

Design	Prompt source	Evidence handling	Conclusion behavior	Risk
A	Pretrained knowledge only	No citations	Uses expected answer	Stale/unsupported
B	Web/search + named sources	Citation after each claim	Marks unsupported claims unverified	Lower grounding risk
C	Old internal summary	Citations to summary only	Uses summary conclusion	May miss recent changes
D	Deep research requested	No source requirements	Assumes retrieved evidence is correct	Evidence may be blended without claim-level checking

- A. A. A
- B. B. B
- C. C. C
- D. D. D

Answer: B