

FDE-Associate Foundation — Advanced Scenario MCQ Practice

Source: ChatGPT & Claude Quick Reference / Claude and ChatGPT 101 (Panaversity AI Agent Factory). Questions are derived primarily from the source's nine concepts and practice material, with additional FDE-style judgment scenarios. The source notes that product features and names change, so durable reasoning patterns matter more than memorizing interface labels.

Distribution: 35 questions total — 19 single-answer, 8 two-answer, and 8 table-based. For two-answer questions, select exactly two options. For all other questions, select exactly one.

Q1. A customer-support team says its AI assistant is “too slow.” A typical request is a two-sentence policy lookup, but 6% of requests involve multi-step exception analysis. The team proposes moving every request to the heaviest available model. What is the best FDE recommendation?

- A. Use the heaviest model for every request because latency is less important than quality.
- B. Route routine lookups to the fast tier and escalate exception cases to thinking/stronger capability based on task complexity.
- C. Keep the current model and add more standing instructions; model choice should never change.
- D. Replace the assistant with a custom GPT so latency becomes predictable.

Answer: B

Q2. An implementation team is deciding whether to memorize current model names or build an internal routing guide. Model names have changed twice during the project. Which design is most durable?

- A. Hard-code vendor model names into every workflow.
- B. Route by task tier—fast, reasoning, flagship—then map those tiers to current vendor names.
- C. Always use the newest model because it supersedes earlier models.
- D. Select models by which vendor has the shortest interface.

Answer: B

Q3. An analyst asks an assistant to summarize a contract but forgets to attach the contract. The assistant produces a polished summary based on generic contract patterns. What is the most important diagnosis?

- A. The assistant needs a heavier reasoning model.
- B. The prompt is too short and needs chain-of-thought instructions.
- C. The task lacks the required source context; better context is the first fix.
- D. The analyst should create a custom GPT before attaching documents.

Answer: C

Q4. A long-running chat has begun contradicting figures established earlier and has forgotten a formatting rule. The team wants to continue the work without losing the useful decisions. What should they do first?

- A. Keep prompting until the model remembers the earlier state.
- B. Ask for a concise state summary, verify it, then start a fresh chat with the needed files/context.
- C. Delete all stored memory immediately.
- D. Switch to the flagship model and assume context degradation disappears.

Answer: B

Q5. A finance team uploads last quarter's financial model into an ongoing project. Three months later, the assistant continues using old numbers when drafting the new board deck. Which failure mode is most relevant?

- A. Model tier mismatch.
- B. Connector permission failure.
- C. Stale persisted context silently influencing outputs.
- D. Research mode overuse.

Answer: C

Q6. A consultant works on the same client every week, repeatedly re-uploading the same policy manuals and restating the audience, tone, and reporting rules. What is the strongest reason to create a project?

- A. A project guarantees factual correctness.

- B. A project makes model reasoning deeper.
- C. The work is a continuing stream with persistent knowledge and standing instructions.
- D. Projects eliminate the need to verify outputs.

Answer: C

Q7. An engineer needs a reusable method for checking every customer contract against the same 14-point checklist, including examples and supporting resources. The customer documents themselves belong to a separate ongoing project. Which abstraction best fits the repeated method?

- A. Memory.
- B. Skill.
- C. Project.
- D. Temporary chat.

Answer: B

Q8. A team says, “Our project and skill are redundant.” They use a customer-preparation project containing current customer files and a skill that specifies how to turn those files into a standardized meeting brief. Which explanation is most accurate?

- A. The project supplies the method; the skill supplies the customer data.
- B. The project supplies scoped knowledge/context; the skill supplies the repeatable method.
- C. Both should be merged because skills cannot use project knowledge.
- D. The skill should store all customer files so the project becomes unnecessary.

Answer: B

Q9. A manager asks an assistant, “What is today’s exchange rate?” The assistant starts a multi-source research process and returns a long cited report. What is the best critique?

- A. Research is always superior when citations are available.
- B. A quick current fact should normally be routed to web search rather than a full research investigation.
- C. Thinking mode should always replace search for current information.
- D. The manager should upload an exchange-rate spreadsheet first.

Answer: B

Q10. A product lead asks: “Compare these three pricing strategies using the assumptions already in this message. No new facts are required.” Which route best matches the work?

- A. Web search.
- B. Research mode.
- C. Thinking/reasoning mode.
- D. Workplace search.

Answer: C

Q11. A compliance team needs a cited report on the current regulatory landscape across five jurisdictions, using multiple external sources. Which route is most appropriate?

- A. Ordinary chat only.
- B. Thinking mode only.
- C. Web search for one fact and stop.
- D. Research mode.

Answer: D

Q12. An FDE connects an AI assistant to a company email system. The assistant can search and read messages and, depending on the tool, may be able to send them. What should the FDE verify before approving the integration?

- A. Only whether the connector has a recognizable brand.
- B. What it can read, what actions it can take, whether the source is trusted, and whether the account is authorized.

C. Whether MCP makes the integration automatically safe.

D. Whether the assistant promises not to make mistakes.

Answer: B

Q13. A team assumes that granting an MCP-connected tool permission to a mailbox automatically gives the assistant every possible action in that mailbox. Which statement corrects the assumption?

A. Permissions can create any missing tool capability.

B. MCP guarantees write access whenever read access is granted.

C. Permissions select among capabilities the tool actually exposes; they do not add new actions.

D. MCP disables all external actions by design.

Answer: C

Q14. A customer asks for a dashboard showing 12 vendors, their prices, and lead times, with the expectation that operations will sort and filter it later. Which output design is strongest?

A. A polished prose paragraph containing all vendors.

B. A structured table or spreadsheet.

C. A long chat response with headings only.

D. A research report without structured fields.

Answer: B

Q15. A CFO will act on a calculated forecast number. The assistant returns a plausible figure and says it “computed” it, but the CFO cannot see which data column was used. What is the best next step?

A. Trust the number because code execution is exact.

B. Ask for the calculation to be performed by code and inspect the code/data column used.

C. Replace the model with a custom GPT.

D. Use research mode to find another forecast online.

Answer: B

Q16. A team compares Claude and ChatGPT on the same internal task. Claude wins on one task; ChatGPT wins on another. The team concludes that one vendor is objectively better. What is the stronger interpretation?

A. The winning assistant should always be used for every task.

B. The comparison demonstrates platform awareness: different tools can have different strengths by task.

C. The result proves benchmarks are unnecessary in all contexts.

D. The team should average the two outputs and discard both originals.

Answer: B

Q17. An FDE tests an assistant on a new production dataset and is impressed because the answer “looks right.” Later, an analyst discovers a subtle calculation error. Which evaluation practice from the source would have reduced this risk earliest?

A. Start with a historical case whose correct result is already known, reproduce it, and compare against the trusted result.

B. Use the most expensive model on the production data.

C. Ask the assistant whether it is confident.

D. Run research mode before every production request.

Answer: A

Q18. An assistant passes five historical test cases exactly. The project owner says, “Great—we can now stop reviewing future outputs.” What is the key correction?

A. Passing the tests transfers accountability to the assistant.

B. Passing earns tested confidence for similar work, but accountability and verification remain with the human.

C. Historical tests are useless unless the assistant fails at least once.

D. The next step is to disable all human review.

Answer: B

Q19. A team is designing a portable workflow that may move between vendors. They want the repeatable method to survive interface changes and avoid being trapped in one vendor's prompt. What is the strongest approach?

- A. Put the method entirely in a vendor-specific chat transcript.
- B. Package the method using the shared Agent Skills approach so it can be reused across compatible tools.
- C. Store the method only in memory.
- D. Encode the method as a temporary chat.

Answer: B

Q20. A healthcare customer wants a scoped workspace for an ongoing clinical-operations project. The team also wants to reduce accidental leakage of unrelated personal context into that workspace. Which TWO choices best match the source's persistence model?

- A. Put stable tone rules into standing instructions.
- B. Put every evolving personal preference into the project instructions.
- C. Use a project for the scoped work and, where available, project-only memory to isolate outside memory.
- D. Put customer files into global memory so every chat can retrieve them.
- E. Treat temporary/incognito chat as permanently trace-free.

Correct answers: A and C

Q21. An FDE is designing a repeated monthly reporting workflow. Which TWO statements correctly distinguish the relevant mechanisms?

- A. Projects are primarily for keeping knowledge/context together for a stream of work.
- B. Skills primarily package a repeatable method for performing a task.
- C. Skills are a replacement for all external connectors.
- D. Projects are intended only for one-off questions.
- E. A custom GPT and a skill are identical abstractions.

Correct answers: A and B

Q22. A customer wants to connect its project-management system to an AI assistant. Before approving the connection, which TWO questions are essential governance checks?

- A. What information can the integration read?
- B. Does the vendor use the newest model family?
- C. What actions can it take externally, such as writing, sending, or deleting?
- D. Does MCP guarantee that the tool is trustworthy?
- E. Can the assistant generate a nicer interface than the existing tool?

Correct answers: A and C

Q23. A research team is unsure whether to use search, thinking, or research. Which TWO mappings are correct?

- A. One current fact → web search.
- B. Hard reasoning over information already provided → thinking/reasoning.
- C. Multi-source current investigation → ordinary chat only.
- D. Deep reasoning with no new information → research mode by default.
- E. Organization-specific information → always public web search.

Correct answers: A and B

Q24. A team is building a financial-analysis assistant. Which TWO practices best align with the source's evaluation discipline?

- A. Test first on a historical case with a trusted correct result.
- B. Let the assistant see the original answer before reproducing it to improve accuracy.
- C. Compare the reproduction with the known result and iterate.
- D. Treat a successful demo as proof for all future data.

- E. Remove human verification after five passing examples.

Correct answers: A and C

Q25. A project has accumulated months of context. The assistant begins contradicting earlier decisions. Which TWO actions are most defensible?

- A. Ask for a concise state summary of decisions, open items, and rules.
- B. Continue prompting indefinitely because context will eventually self-correct.
- C. Verify the summary before carrying it into a fresh conversation.
- D. Delete the project immediately without preserving useful state.
- E. Assume the problem is always model intelligence rather than context.

Correct answers: A and C

Q26. An operations team needs a reusable output for a recurring workflow. Which TWO design choices improve downstream usability?

- A. Specify the finished object's behavior and audience rather than only naming it.
- B. Ask for operational data as structured tables/spreadsheets when it will be sorted or loaded elsewhere.
- C. Prefer prose for data because prose gives the model more freedom.
- D. Avoid specifying format so the model can choose.
- E. Put every revision into one unstructured chat message.

Correct answers: A and B

Q27. A multinational organization is evaluating Claude and ChatGPT. Which TWO principles from the source should guide the decision?

- A. Learn the transferable workflow discipline and then learn where each product puts its controls.
- B. Choose the platform solely by interface familiarity.
- C. Compare the same task across assistants when the work matters.
- D. Assume model names are stable enough to be the primary mental model.
- E. Treat identical control names as evidence that the products behave identically.

Correct answers: A and C

Q28. A logistics company is deciding how to route AI work. The FDE maps each request to the least costly route that still fits the job. Based on the table, which routing decision is most appropriate for Request 3?

Request	Information state	Work type	Preferred route	Why
1	One current public fact	Quick lookup	Web search	Fast current answer
2	All facts already supplied	Complex comparison	Thinking	Reasoning, no new info
3	Current market across 6 countries	Cited multi-source investigation	Research	Plans and investigates
4	Internal approved policy	Organization knowledge	Workplace search/app	Uses connected internal systems

Question: Select the best answer based on the table.

- A. Web search
- B. Thinking
- C. Research
- D. Workplace search/app

Answer: C

Q29. A software team is deciding where to place different kinds of persistent context. Which row is internally consistent with the source?

Item	Need	Best home	Reason	Risk if misplaced
A	Stable response style	Standing instructions	Broad stable rule	Project-specific scope
B	Current customer files	Project	Scoped knowledge	Global persistence
C	Repeatable contract-review method	Skill	Packages method	Re-uploading instructions

Item	Need	Best home	Reason	Risk if misplaced
D	Evolving personal preference	Project-only memory	Scoped work context	Loss of global preference

Question: Select the best answer based on the table.

- A. Row A only
- B. Rows A, B and C
- C. Rows B and D
- D. All rows

Answer: B

Q30. An FDE audits an integration before production. The tool is MCP-based. The table records what the tool exposes and what the customer account permits. Which conclusion is best supported?

Capability	Tool exposes?	Account permits?	Result	FDE implication
Read email	Yes	Yes	Available	Can search/read
Send email	No	Yes	Unavailable	Permission cannot create capability
Delete email	Yes	No	Unavailable	Capability exists but access is denied
Create calendar event	Yes	Yes	Available	Can act if workflow authorizes it

Question: Select the best answer based on the table.

- A. Send email is available because the account permits it.
- B. Delete email is available because the tool exposes it.
- C. The integration's actions are the intersection of exposed capabilities and granted permissions.
- D. MCP makes every exposed action automatically available.

Answer: C

Q31. A board-reporting project contains persistent context. The FDE observes the following evidence. Which row identifies the most likely stale-context failure?

Observation	Likely cause	Best first response	Confidence
A	Model too weak	Switch to flagship	Low
B	Old financial model still stored	Review/update project knowledge	High
C	Connector missing	Install connector	Low
D	Research too shallow	Run research	Low

Question: Select the best answer based on the table.

- A. A
- B. B
- C. C
- D. D

Answer: B

Q32. A customer wants to turn an AI-generated deliverable into something operations can actually use. Which row best matches the output design principle?

Deliverable	Downstream need	Recommended form	Reason	Poor alternative
1	Read a narrative memo	Paragraph/document	Human reading	Spreadsheet
2	Sort/filter vendor records	Table/spreadsheet	Structured fields	Paragraph
3	Act on a computed forecast	Code-computed result	Check calculation path	Unverified prose number
4	Revise a working mini-app	Artifact/writing/code work product	Separate editable object	Long chat scroll

Question: Select the best answer based on the table.

- A. Deliverable 1 only
- B. Deliverables 2 and 3 only

C. Deliverables 2, 3 and 4

D. All four require a spreadsheet

Answer: C

Q33. A team compares assistants using known historical cases. The table shows the evaluation results. What conclusion is justified?

Case	Known correct result?	Assistant result	Matched?	Implication
1	Yes	Exact match	Yes	Evidence of capability
2	Yes	Missed key pattern	No	Needs refinement
3	Yes	Exact match	Yes	Evidence of capability
4	No	Plausible answer	Unknown	Not a validated pass

Question: Select the best answer based on the table.

A. Four cases passed because all outputs were plausible.

B. Three cases provide validated evidence and one remains unvalidated.

C. Case 2 should be ignored because most cases passed.

D. Case 4 is stronger evidence than Case 1 because it used new data.

Answer: B

Q34. An FDE is deciding whether to use a custom GPT, project, or skill for three needs. Based on the table, which mapping is best?

Need	Primary requirement	Best mechanism	Why
1	Keep one client's files/chats/instructions together	Project	Scoped context
2	Package a recurring checklist and examples	Skill	Repeatable method
3	Create a separately configured assistant with its own persona/tools	Custom GPT	Distinct assistant configuration
4	Connect to an external system	Project	Projects are the integration layer

Question: Select the best answer based on the table.

A. 1→Project, 2→Skill, 3→Custom GPT

B. 1→Skill, 2→Project, 3→Custom GPT

C. 1→Project, 2→Custom GPT, 3→Skill

D. 1→Custom GPT, 2→Skill, 3→Project

Answer: A

Q35. A customer says a temporary/incognito conversation is safe enough for highly sensitive work because “nothing is retained.” The FDE reviews the source’s stated behavior. Which conclusion is most accurate?

Mode	Normal history?	Normal memory?	Limited vendor retention?	Correct interpretation
A	Yes	Yes	Possible	Persistent
B	No	No	Yes, for safety/abuse review	Memory-free, not trace-free
C	No	No	No	Completely unretained
D	Yes	No	Yes	Memory-free but historical

Question: Select the best answer based on the table.

A. A

B. B

C. C

D. D

Answer: B