

FDE-Associate Foundation — Workflow Design & Diagnosis

35 scenario-based MCQs • 19 single-answer • 8 two-answer • 8 table-based

Source reviewed: Panaversity / Agent Factory — “Workflow Design & Diagnosis” Crash Course. Exam emphasizes delegation maps, requirements, computation, window tests, diagnostic timing, durable fixes, measurement, incident response, and accurate communication.

Instructions: Questions are intentionally difficult. For checkbox questions, select exactly two answers. Table questions require reasoning from the supplied table.

1. A procurement team wants AI to review vendor contracts. The AI can reliably extract clauses and draft deviations, but the final approval creates binding obligations. The project lead proposes giving AI approval because the drafting accuracy is 98% on a pilot. Which decision best follows the course's delegation logic?

- A. Delegate approval because the pilot establishes end-to-end reliability.
- B. Keep approval human because drafting quality is evidence about drafting, not the decision step.
- C. Delegate approval but add a stronger prompt to force caution.
- D. Delegate approval only for clauses that look mechanically simple.

Answer: B

2. A 40-page policy, several emails, and meeting notes must become requirements for an AI-assisted workflow. Which acceptance test most strongly indicates that a requirement is actually settled?

- A. Everyone agrees that the wording sounds reasonable.
- B. The requirement contains at least one AI-related term.
- C. You can state what evidence would prove the requirement was met.
- D. The requirement is short enough to fit in a prompt.

Answer: C

3. A finance workflow asks an AI assistant to report quarterly revenue growth. The model produces a plausible percentage in prose, but the source spreadsheet has several possible revenue columns. What is the best first redesign?

- A. Ask the model to be more careful with arithmetic.
- B. Increase the model tier.
- C. Use code execution to compute the figure and verify which column was used.
- D. Ask two models for the percentage and average their answers.

Answer: C

4. An operations plan recommends hiring three analysts from ticket trends. The arithmetic is correct, but an internal hiring freeze was decided in a meeting and is not in the supplied data. Which diagnosis is strongest?

- A. The calculation step is broken.
- B. The synthesis step failed because models cannot plan staffing.
- C. The judgment step lacks an input that cannot be supplied from the available context.
- D. The prompt needs more formatting constraints.

Answer: C

5. A team maps an onboarding workflow and classifies each step using three criteria. Which combination is the intended basis for deciding ownership?

- A. Model confidence, frequency, and token cost.
- B. Reversibility, stakes, and accountability.
- C. Complexity, novelty, and user preference.
- D. Accuracy, latency, and context length.

Answer: B

6. An AI workflow produces a correct draft every Monday. The team then lets the AI make the next decision without review because the draft has been consistently strong. What failure pattern is this?

- A. Context overload.
- B. Stale configuration.
- C. Halo delegation.
- D. Under-specification.

Answer: C

7. A collaborative review gate is still present on the workflow diagram, but reviewers are overloaded and routinely approve AI drafts with only a glance. What is the most accurate diagnosis?

- A. The map is necessarily wrong because every collaborative step must be removed.
- B. The gate is unstaffed in practice, even though it remains on the map.
- C. The model has become less capable.
- D. The workflow has become under-specified.

Answer: B

8. A team designs a workflow around a highly capable summarization Skill. They add steps mainly because the Skill supports them, rather than because the underlying business process requires them. What is the core mapping error?

- A. Over-delegation.
- B. Halo delegation.
- C. Mapping the tool instead of the work.
- D. Context overload.

Answer: C

9. A prompt works perfectly for the first ten turns of a long investigation, then the assistant starts ignoring an output-format requirement that was established early. What is the strongest initial hypothesis?

- A. Stale configuration.
- B. Context overload.
- C. Expectation mismatch.
- D. Wrong accountability assignment.

Answer: B

10. A workflow's output has been wrong from the very first run because the prompt never states what counts as a 'key issue.' Which diagnosis best fits?

- A. Under-specification.
- B. Context overload.
- C. Stale configuration.
- D. Wrong model tier.

Answer: A

11. A workflow used to produce correct reports, but after several months the reports begin quoting last quarter's figures while formatting remains correct and no errors are raised. What should be checked first?

- A. Whether the model can summarize.
- B. Whether the conversation is too long.
- C. Whether the configuration or knowledge source is stale.
- D. Whether the task should be abandoned.

Answer: C

12. A user reports that a model 'cannot do the task' because the output is poor. According to the course's diagnostic sequence, what should happen before questioning task fit?

- A. Immediately move to the strongest model.
- B. Re-check prompt, session length, feature/model choice, and configuration.
- C. Rebuild the entire workflow.
- D. Add a human review gate regardless of cause.

Answer: B

13. A recurring correction is: 'Always include the customer's account ID in the report.' It is a stable rule, not a multi-step procedure. Where should the correction be promoted?

- A. A standing instruction.
- B. A knowledge-base reference file.
- C. A Skill containing the full procedure.
- D. The incident log only.

Answer: A

14. A recurring correction says: 'First validate the input schema, then normalize fields, then generate the report.' It has been fixed manually several times. What is the best durable home?

- A. Standing instruction only.
- B. Knowledge-base reference only.
- C. Skill/procedure.
- D. A longer user prompt each time.

Answer: C

15. A team wants to determine whether an AI redesign actually improved a weekly process. Which measurement approach best follows the course?

- A. Compare the redesigned run with memory of how the old process felt.
- B. Establish a baseline before changing the workflow, then compare time, revisions, and manual steps.
- C. Measure only model response latency.
- D. Ask the AI whether the redesign is better.

Answer: B

16. A customer complaint workflow sends a message externally. The message can be generated by AI, but once sent it cannot realistically be undone. Who should own the sending step?

- A. AI, because sending is mechanical.
- B. AI, if complaint severity is low.
- C. A named human owner because the step is irreversible and externally binding.
- D. Whoever wrote the prompt.

Answer: C

17. A team wants to redesign a workflow but has not yet identified which numbers matter. What is the best design move?

- A. Start building the agent and discover metrics later.
- B. Extract requirements and explicitly identify numbers that must be computed.
- C. Ask the model for a polished architecture first.
- D. Choose the fastest model and optimize after deployment.

Answer: B

18. A manager asks, 'Does the system automate contract review?' The actual workflow has AI extraction and redlining, followed by mandatory lawyer approval of every change. Which answer is most defensible?

- A. Yes, it fully automates contract review.
- B. Yes, AI handles contract review with occasional oversight.
- C. It assists extraction and drafting; a lawyer still reviews and approves every change before release.
- D. It is basically as good as a lawyer for standard contracts.

Answer: C

19. A workflow designer is tempted to move a task to AI because it is repetitive and easy to describe. The task also determines whether a customer receives a legally binding commitment. What should dominate the ownership decision?

- A. Repetition alone.
- B. The tool's demonstrated drafting accuracy.
- C. Reversibility, stakes, and who is accountable.
- D. Whether the task can fit in one prompt.

Answer: C

20. A team is redesigning a weekly customer-refund workflow. Which two findings should push the designer toward retaining a human owner for specific steps? Select exactly two.

- A. The step can be undone immediately with no external effect.
- B. The step creates an externally binding commitment.
- C. The model has handled a neighboring drafting step accurately.
- D. The step's decision depends on an internal risk appetite that is not in the supplied context.
- E. The step has a high token count.
- F. The step is repetitive.

Correct answers: B and D

21. A workflow is producing poor results. Which two observations are especially diagnostic of context overload rather than under-specification? Select exactly two.

- A. The first answer was wrong because the prompt omitted a required criterion.
- B. The output was good initially but degraded as the session became long.
- C. A fresh session restores compliance with the earlier instruction.
- D. The same error occurs on the first answer every time.
- E. The knowledge source has not been updated for a year.
- F. The task requires a legal judgment.

Correct answers: B and C

22. A team wants to make a recurring fix durable. Which two pairings are correct? Select exactly two.

- A. Stable rule → standing instruction.

- B. Factual reference material → knowledge base.
- C. Multi-step procedure → one sentence in the prompt.
- D. Multi-step procedure → knowledge base only.
- E. Stable rule → incident log only.
- F. Reference material → model tier upgrade.

Correct answers: A and B

23. A customer asks what an AI workflow actually does. Which two communication choices best preserve trust? Select exactly two.

- A. Say 'fully automated' if the model is accurate on the pilot.
- B. State what the AI does and explicitly name the human checkpoint.
- C. Say 'AI handles X' and assume the audience knows review exists.
- D. State important refusal/boundary cases where the workflow hands work to a person.
- E. Say it is 'as good as a person' for normal cases.
- F. Avoid describing limitations unless asked.

Correct answers: B and D

24. A team discovers a bad output escaped a workflow. Which two actions belong in the first-hour response described by the course? Select exactly two.

- A. Immediately tune the prompt before determining scope.
- B. Stop the spread using a known halt mechanism.
- C. Wait until the next scheduled map review to document it.
- D. Determine how many outputs were affected and where they went.
- E. Quietly delete the bad output and continue.
- F. Replace the model before recording facts.

Correct answers: B and D

25. A workflow redesign is being evaluated. Which two practices provide the strongest evidence that the change is worth keeping? Select exactly two.

- A. Establish a baseline before the change.
- B. Measure only whether the AI's prose sounds better.
- C. Track time, revision rounds, and manual steps.
- D. Skip measurement if the team likes the new workflow.
- E. Let the model self-certify the improvement.
- F. Change multiple unrelated variables at once.

Correct answers: A and C

26. A task may or may not fit the AI workflow. Which two statements correctly reflect the course's 'window test' logic? Select exactly two.

- A. A step is suitable when every input it needs can be put in the model's context.
- B. Any judgment step must remain human because models cannot perform judgment.
- C. A missing organizational decision in a meeting can make an otherwise well-specified step unsuitable for AI ownership.
- D. If the model can guess a missing input, the window test passes.
- E. The window test is primarily about token price.
- F. The context window is irrelevant if the prompt is precise.

Correct answers: A and C

27. A team is diagnosing four possible causes of bad output. Which two practices follow the intended diagnostic discipline? Select exactly two.

- A. Start by switching to the most capable model.
- B. Check the prompt against role, context, task, constraints, and format.
- C. Check session length before changing configuration.
- D. Declare the task impossible after the first bad result.
- E. Skip configuration checks because stale configuration always throws errors.
- F. Diagnose only from the quality of the final prose.

Correct answers: B and C

28. A healthcare operations team is piloting an AI-assisted referral workflow. The designer must assign ownership before implementation. The table records the nature of each step and what happens if it is wrong. The team wants the highest-risk irreversible action to remain accountable to a person while allowing reversible work to be delegated.

Step	Reversible?	Bad-case consequence	Missing/available inputs	Candidate owner
Extract referral fields	Yes	Rework only	All source fields available	AI
Draft referral summary	Yes	Reviewer revises	All clinical notes supplied	AI
Decide escalation category	No	Urgent case may be missed	Policy supplied; local risk appetite absent	AI / Human
Send escalation notice	No	External action and patient impact	Approved category available	Human
Archive final record	Mostly	Audit trail affected	All fields available	AI / Human

Question: Based on the table, which ownership decision is the strongest application of reversibility, stakes, and accountability?

- A. AI should own escalation categorization because policy is supplied.
- B. Human should own escalation categorization because a critical input—local risk appetite—is unavailable and the consequence is irreversible.
- C. AI should own sending because the category is already approved.
- D. Human should own all five steps because healthcare is high stakes.

Answer: B

29. A finance team is troubleshooting a monthly reporting assistant. The team records when each failure begins and what happens in a fresh session or with a different feature. Use the table to select the best diagnosis for the case that should be addressed before declaring the task a poor fit.

Case	First symptom	Fresh session	Different feature/tier	Configuration status
A	Wrong on first answer	Still wrong	Still wrong	Prompt omits required criterion
B	Correct for 12 turns, then format degrades	Format returns	Not tested	Current
C	Arithmetic wrong in repeatable way	Still wrong	Code execution corrects it	Current
D	Worked last quarter, now stale figures	Still wrong	Still wrong	Knowledge source outdated

Question: Which case most clearly points to under-specification and therefore warrants a prompt/requirements fix before any model change?

- A. Case A
- B. Case B
- C. Case C
- D. Case D

Answer: A

30. A sales-operations team has converted a messy RFP process into explicit requirements. The designer must place each reusable component in the correct home so future projects inherit the right behavior without duplicating setup. The table shows the nature of each item.

Item	Nature	Changes by project?	Needs ordered steps?	Best home
Past winning proposals	Reference material	Yes / project-specific	No	Project knowledge base
'Always return requirements in JSON'	Stable rule	No	No	Standing instruction

Item	Nature	Changes by project?	Needs ordered steps?	Best home
Validate schema → normalize → generate	Procedure	No	Yes	Skill
Quarterly sales figures	Reference data	Yes	No	Project knowledge base

Question: Which mapping is correct and most durable according to the course?

- A. Past proposals → Skill; stable JSON rule → knowledge base; procedure → standing instruction.
- B. Past proposals → standing instruction; stable JSON rule → Skill; procedure → knowledge base.
- C. Past proposals → project knowledge base; stable JSON rule → standing instruction; ordered procedure → Skill.
- D. All three should be embedded in each user's prompt.

Answer: C

31. A customer-support workflow has been operating for three months. Leadership wants to know whether the map still represents reality. The table contains evidence about designed ownership and actual behavior.

Step	Designed owner	Actual behavior	Risk signal	What changed?
Classify severity	AI	AI	High: classification controls routing	Nothing
Draft response	Collaborative	Reviewer reads every case	Medium	Volume doubled
Send response	Human	Low-severity cases bypass review	High: irreversible	A rule was added outside the map
Log outcome	AI	AI	Low/reversible	Nothing

Question: Which row is the clearest evidence that the designed map no longer matches the actual workflow and requires immediate attention?

- A. Classify severity
- B. Draft response
- C. Send response
- D. Log outcome

Answer: C

32. A product team is deciding whether to keep an AI workflow after a pilot. The table compares a pre-change baseline with two candidate redesigns. Use the data rather than intuition.

Metric	Baseline	Redesign A	Redesign B
Average handling time	42 min	29 min	31 min
Revision rounds	3.1	2.8	1.4
Manual steps	8	6	3
Bad outputs reaching customers	0.2%	0.9%	0.2%
Measurement quality	Measured	Measured	Estimated

Question: Which redesign is the strongest candidate for controlled adoption based on the evidence shown?

- A. Redesign A, because it has the lowest handling time.
- B. Redesign B, because it reduces manual steps and revision rounds while keeping bad-output rate at baseline.
- C. Redesign A, because its measurements are cleaner than B's.
- D. Redesign B, because estimated measurements are sufficient for production decisions.

Answer: B

33. A legal team needs to explain an AI-assisted contract workflow to different audiences. The table shows each audience's primary concern and candidate wording. Select the wording that best matches the audience's need while preserving the human gate.

Audience	Primary concern	Candidate wording	Gate explicit?
Legal lead	Mechanism and failure modes	AI extracts clauses and drafts redlines; counsel reviews every change and watches for implied obligations.	Yes
Executive	Outcome and oversight	The process is faster, with counsel still approving every change before release.	Yes
Risk function	Control	Drafting is AI-assisted; a qualified lawyer reviews and approves every term before sending.	Yes
Client	Simple benefit	AI fully automates standard contract review.	No

Question: Which audience-wording pairing best satisfies the course's communication principle?

- A. Client → 'AI fully automates standard contract review.'
- B. Executive → 'AI extracts clauses and watches for implied obligations.'
- C. Risk function → 'Drafting is AI-assisted; a qualified lawyer reviews and approves every term before sending.'
- D. Legal lead → 'The process is faster, with counsel still approving changes.'

Answer: C

34. A manufacturing team sees several recurring defects in an AI work-order assistant. The table captures symptom timing and likely diagnosis. The engineer must choose the case where changing the model tier would be the least justified first move.

Case	Timing	Pattern	Current config	Best first check
A	Wrong from first run	Missing safety criterion	Current	Prompt completeness
B	Degrades after long session	Format drift	Current	Fresh session / summary
C	One repeatable numeric error	Same column misread	Current	Code execution / feature
D	Worked reliably, then deteriorated	Stale maintenance data	Old	Configuration/knowledge date

Question: Which case should NOT trigger an immediate model-tier upgrade because a cheaper diagnostic is indicated by the timing/pattern evidence?

- A. A
- B. B
- C. C
- D. D

Answer: A

35. A company has just experienced a bad output from a customer-facing workflow. The incident commander reviews the first-hour plan. The table shows possible actions and whether they address the required incident dimensions.

Action	Stops spread?	Scopes impact?	Notifies stakeholders?	Changes workflow map?
A	Yes	No	No	No

Action	Stops spread?	Scopes impact?	Notifies stakeholders?	Changes workflow map?
B	No	Yes	No	No
C	No	Yes	Yes	No
D	No	No	Yes	Yes

Question: Which action is the only one in the table that both stops further spread and therefore satisfies the first-hour containment requirement?

- A. A
- B. B
- C. C
- D. D

Answer: A

Exam Blueprint Verification

Count check completed before export:

- Single-answer: 19 questions (Q1–Q19)
- Two-answer checkbox: 8 questions (Q20–Q27)
- Table-based: 8 questions (Q28–Q35)
- Total: 35 questions

Core source principles reflected in the exam include: the step as the unit of value; requirements that can be tested; consequence-based computation; the window test; ownership by reversibility/stakes/accountability; halo delegation and unstaffed gates; diagnosis by symptom timing; cheapest-first troubleshooting; promotion of fixes into rules/references/procedures; baseline measurement; first-hour incident response; and explicit human gates in communication.