

FDE Associate Foundation — Workflow Design & Diagnosis

35 Scenario-Based MCQs | 19 single-answer • 8 two-answer • 8 table-based

Source reviewed: Workflow Design & Diagnosis — The AI Agent Factory / Panaversity. Questions emphasize delegation maps, requirements, build loops, diagnostic timing, persistent fixes, friction, measurement, incident response, and accurate communication. The exam intentionally adds realistic FDE implementation scenarios beyond direct recall.

1. A regional insurer wants AI to draft claim summaries. The draft is reversible, but the final severity code determines whether a claim is escalated and can affect statutory reporting. The operations lead says, “The model has been right for months, so let it assign low/medium/high too.” What is the strongest FDE response?

- A. Keep severity classification human because accountability and downstream consequences differ from drafting.
- B. Automate severity only for claims where the draft confidence exceeds a fixed threshold.
- C. Automate severity because repeated drafting accuracy is evidence that the model understands the case.
- D. Remove the severity field and let the downstream team infer it from the generated summary.

Answer: A

2. A customer says, “We need a weekly report delivered promptly after month-end.” Two teams interpret this differently: one targets two business days and the other targets ten. You are asked to redesign the workflow before implementation begins. What should you do first?

- A. Pick two business days because it is the more customer-friendly interpretation.
- B. Convert the statement into a task definition covering output, audience, frequency, data, and format, then surface the missing deadline explicitly.
- C. Ask the model to infer the deadline from comparable finance workflows.
- D. Prototype both schedules and let the customer choose after seeing the outputs.

Answer: B

3. A finance workflow asks an AI assistant to produce a monthly revenue total from transaction rows. The output looks plausible, and reviewers usually accept it because the number is within the expected range. During an audit, one total is found to be subtly wrong. Which redesign best addresses the failure?

- A. Increase the prompt length so the model has more context about revenue.
- B. Use a stronger model tier while keeping the calculation in prose.
- C. Move the arithmetic into code execution and have the AI explain or present the computed result.
- D. Ask the reviewer to inspect every digit manually before approving the report.

Answer: C

4. A customer-success team has a complaint workflow: retrieve ticket history, summarize the complaint, classify severity, draft a response, send it, and log the outcome. The summary has been accurate for three months, but a missed high-severity complaint would not be recoverable. Which mapping is most defensible?

- A. Let AI classify severity automatically because the summary is already accurate.
- B. Let AI classify severity and require review only when the confidence score is low.
- C. Let AI send low-severity responses automatically because sending is routine.
- D. Treat severity classification as at least collaborative and retain human ownership of the irreversible send decision.

Answer: D

5. A product team is building an internal dashboard through an AI artifact workflow. The first prototype is useful, but users then request role-based access, uptime expectations, auditability, and a guarantee that board-report numbers remain stable across Mondays. What is the key signal?

- A. Other people now depend on the artifact, so it has crossed into infrastructure and needs engineering ownership.
- B. The prompt is becoming too long, so the team should start a new chat for every department.
- C. The artifact should be abandoned because AI prototypes are unsuitable for internal tools.
- D. The team should add more screenshots to the prompt and continue iterating informally.

Answer: A

6. A workflow worked well during a pilot. After an hour-long working session, the assistant starts ignoring a formatting rule that it followed at the beginning. A fresh session using the same initial instruction follows the rule again. What is the best first diagnosis?

- A. Stale configuration, because the original workflow has probably changed.
- B. Context overload, because the symptom appeared after the session accumulated substantial context.
- C. Expectation mismatch, because formatting is outside the model's capabilities.
- D. Wrong feature/model, because any formatting drift requires a stronger model.

Answer: B

7. A reporting assistant has produced correct outputs for several months. This week it begins quoting last quarter's product definitions even though the prompt has not changed. No error is raised, and the reports remain well formatted. What should you investigate first?

- A. Rewrite the prompt to emphasize current-quarter terminology.
- B. Switch immediately to the most capable model tier.
- C. Audit the saved instructions, knowledge sources, and Skills for stale configuration.
- D. Conclude that the task is fundamentally unreliable and return it to a manual process.

Answer: C

8. An analyst repeatedly fixes the same output issue: the executive summary omits the target audience in its first line. The correction is needed for every person using the shared workflow. Where should the permanent fix live?

- A. In each user's memory, because the assistant can learn the preference over time.
- B. In the source data, by adding a new field called "audience."
- C. In a Skill, because every correction is automatically a procedure.
- D. In a standing instruction, because the missing first-line requirement is a rule that should apply to every run.

Answer: D

9. A team notices three kinds of friction in a weekly reporting process: everyone re-pastes the same background, each analyst reformats the report, and different analysts catch different defects. You have one optimization cycle available. Which sequence best follows the course's reasoning?

- A. Consolidate repeated context first, then promote the remaining reference/procedure/rule fixes into shared configuration.
- B. Create three Skills immediately, one for each friction, before reducing the number of workflow steps.
- C. Replace the model before changing workflow structure because variance implies insufficient model capability.
- D. Measure only elapsed time and optimize whichever individual step is longest.

Answer: A

10. A compliance workflow is being redesigned. Management says the goal is "make it faster," but one wrong figure could trigger a material reporting problem. The team proposes optimizing for minutes saved because that is easiest to measure. What is the strongest correction?

- A. Optimize for time because every compliance workflow should minimize human review.
- B. Optimize for accuracy because the business consequence of a wrong figure dominates the value of time saved.
- C. Optimize for variance because multiple analysts are involved.
- D. Avoid metrics because quantitative targets can encourage the model to game the workflow.

Answer: B

11. A team changed its customer-reporting workflow and claims it reduced cycle time from 45 minutes to 25. No one recorded an unchanged run before the redesign; the 45-minute number came from memory. The manager asks how confident they should be in the claimed improvement. What is the best answer?

- A. The improvement is valid because the current process is visibly faster.
- B. The improvement is valid if three analysts agree that 45 minutes feels representative.
- C. The claim is weak because the baseline was not measured before the change and cannot be reconstructed reliably afterward.
- D. The team should estimate the baseline using the new workflow's logs and back-calculate the old time.

Answer: C

12. A new AI workflow is ready to replace an existing manual process used by three analysts. The team wants evidence that the new process is better and a fallback if a defect appears. What is the most appropriate rollout?

- A. Switch immediately because parallel operation doubles cost and therefore proves the new workflow is inefficient.
- B. Let one analyst use the new process permanently while the others keep the old process.
- C. Run the new process only on unusual cases so its limitations are easier to observe.
- D. Run the old and new workflows in parallel for a few cycles, compare the agreed metric, then switch when the evidence is clear.

Answer: D

13. A workflow has a collaborative review gate after an AI-generated draft. The map says a reviewer must inspect every draft, but volume has doubled and the reviewer now approves most items in seconds without opening the source. No incident has happened yet. What is the most accurate diagnosis?

- A. The gate is becoming unstaffed in practice even though the documented design is unchanged.
- B. The gate is still valid because the reviewer technically clicks approve.
- C. The AI step should be made human-retained because all collaborative gates eventually fail.
- D. The map is correct until a customer complains, so no action is justified.

Answer: A

14. A delegation map lists an owner as "Operations Team," has no next review date, and contains a collaborative step without naming who reviews it or what they check. The workflow itself is currently performing acceptably. Which change most directly turns the map into an operating document?

- A. Add more technical details about the model and its temperature.
- B. Name one accountable map owner, set a review date, and specify the reviewer, timing, and check for each gate.
- C. Remove all human checkpoints so the workflow has fewer dependencies.
- D. Replace the map with a prompt that reproduces the current workflow.

Answer: B

15. A customer reports that a generated summary was wrong from the very first response. The prompt contains the task but omits the business definition of "key points," and the model had no access to that definition elsewhere. Which diagnostic step is most justified before changing models?

- A. Increase the context window by moving to a larger model.
- B. Inspect the knowledge base for stale files.
- C. Re-read the prompt against role, context, task, constraints, and format, then add the missing specification.
- D. Decide the task is an expectation mismatch because the model cannot know business priorities.

Answer: C

16. An FDE is asked to describe an AI-assisted contract workflow to a risk committee. The AI extracts clauses, flags deviations, and drafts redlines, but a qualified lawyer approves every change before anything leaves the organization. Which description is most defensible?

- A. "The AI handles contract review with minimal human involvement."
- B. "The workflow is fully automated for standard contracts."
- C. "The model is basically as good as a lawyer for routine clauses."
- D. "AI assists with extraction, deviation flags, and drafting; a qualified lawyer reviews and approves every change before release."

Answer: D

17. A bad output has escaped a customer-facing workflow. The team has paused the process, but nobody knows how many prior outputs may be affected because the configuration became stale sometime during the month. What should the incident response prioritize after stopping further spread?

- A. Determine scope by identifying how many outputs were produced under the suspect configuration and where they went.
- B. Rewrite the prompt immediately so the next run demonstrates that the problem is fixed.
- C. Delete the stale configuration and rely on the latest model to regenerate all outputs.
- D. Wait for affected customers to report problems before estimating scope.

Answer: A

18. A sales forecasting team asks the assistant to predict the exact revenue number for next quarter. The prompt is detailed, the model is strong, the context is current, and the workflow has been stable. They still want one exact number that will be guaranteed correct. What is the best FDE judgment?

- A. Add more historical examples until the exact value becomes predictable.
- B. Treat this as an expectation mismatch and reshape the task around a range, assumptions, or adjustable drivers.
- C. Move the request into a Skill so the procedure becomes deterministic.
- D. Add a human approval gate and then keep the exact forecast requirement unchanged.

Answer: B

19. A workflow has five steps. The AI drafts a response, a named support lead judges it, and then a system sends the approved response externally. The support lead says, "I review when I have time." Which interpretation is strongest?

- A. The workflow is collaborative because a reviewer exists somewhere in the organization.
- B. The workflow is human-retained because sending is external.
- C. The collaborative gate is not operationally defined until the reviewer and moment/check are explicit; otherwise it effectively behaves as automated.
- D. The gate can be omitted because the AI draft is reversible before sending.

Answer: C

20. A manufacturing customer wants AI to inspect work orders and produce a dispatch recommendation. The final dispatch changes a physical schedule and can create costly downstream commitments. Which TWO design decisions best align with the source's delegation criteria?

- A. Let AI own dispatch because the recommendation is usually correct.
- B. Let AI extract and normalize work-order data where errors can be checked downstream.
- C. Let AI execute the final dispatch when confidence exceeds 95%.
- D. Keep the final dispatch decision with a named accountable person.
- E. Infer dispatch authority from the quality of the preceding extraction step.
- F. Treat all steps as human because physical operations are inherently too risky.

Correct answers: B and D

21. A team is converting a 30-page policy and several emails into an implementable workflow. The request contains vague phrases such as “reasonable turnaround” and “appropriate escalation.” Which TWO actions are strongest before implementation?

- A. Ask the model to choose likely defaults so engineering can keep moving.
- B. Extract distinct requirements with source locations and mark ambiguous or implied requirements.
- C. Turn the entire policy into one long standing instruction.
- D. Resolve the undefined terms into testable criteria with the accountable stakeholders.
- E. Start building the prototype and discover the definitions through user feedback.
- F. Replace ambiguous requirements with the model's most common interpretation.

Correct answers: B and D

22. An AI dashboard prototype is useful, but a year later multiple departments depend on it for weekly management reporting. Which TWO signals justify escalation beyond prompt-and-iterate?

- A. Other people now depend on it as infrastructure.
- B. The artifact contains more than five charts.
- C. It now needs uptime, access control, ownership, and repeatability guarantees.
- D. The team has used the same prompt more than three times.
- E. Users have asked for a new color scheme.
- F. The prototype was originally created in a chat interface.

Correct answers: A and C

23. A support workflow's output starts failing after long sessions, but the same initial request works correctly in a fresh session. Which TWO responses are most appropriate?

- A. Treat the timing as evidence for context overload.
- B. Replace the model before checking session length.
- C. Restart from a concise summary or otherwise reduce accumulated context.
- D. Declare the task an expectation mismatch because long conversations are inherently unsafe.
- E. Add more unrelated examples to the current session to reinforce the instruction.
- F. Update the knowledge base before testing a fresh-session hypothesis.

Correct answers: A and C

24. An analyst has repeatedly corrected a shared report workflow. The correction says the report must always exclude canceled orders and use the current product catalog. Which TWO permanent homes are correct?

- A. Put “exclude canceled orders” in a standing instruction.
- B. Put the current product catalog in the knowledge base.
- C. Put both corrections into one Skill regardless of their nature.
- D. Leave both corrections in chat because they are analyst judgment.
- E. Store the catalog as a standing instruction so it is always visible.

■ F. Put the exclusion rule into the knowledge base as reference material.

Correct answers: A and B

25. A team wants to optimize a customer-facing reporting workflow. Different analysts produce slightly different formats, and reviewers spend time reconciling them. Which TWO actions best address the actual friction?

- A. Optimize only elapsed time because consistency is subjective.
- B. Establish a shared procedure/Skill for the required report format.
- C. Give each analyst a larger context window so their personal style is preserved.
- D. Treat cross-person variance as a workflow-level problem rather than an individual performance issue.
- E. Add more optional output formats so analysts can choose the best one.
- F. Remove the reviewer because reconciliation is evidence that review is redundant.

Correct answers: B and D

26. A workflow sends an incorrect customer notice. The team has already stopped further sends. Which TWO actions belong in the designed first-hour response?

- A. Determine how many notices were affected and where they went.
- B. Wait until the root cause is proven before telling anyone.
- C. Record the facts and follow the organization's reporting path, including the appropriate owner/risk function.
- D. Immediately rewrite every prompt without preserving evidence.
- E. Keep the incident out of the workflow map because the map should describe the ideal state.
- F. Restore the workflow immediately so the team can compare old and new outputs.

Correct answers: A and C

27. A workflow has a repeated manual correction that is clearly a multi-step sequence: retrieve the latest data, validate two fields, transform the data in a fixed order, then produce a standardized artifact. Which TWO conclusions are strongest?

- A. It is probably a procedure rather than a one-line rule.
- B. It belongs in a standing instruction because every repeated fix is a rule.
- C. A Skill is an appropriate home if the sequence is stable and shared.
- D. It should remain in individual chat history to preserve flexibility.
- E. The knowledge base is the best home because procedures are reference material.
- F. Memory is preferable because the assistant can learn the sequence implicitly.

Correct answers: A and C

28. A healthcare-adjacent benefits team is mapping an eligibility workflow. The team wants to maximize automation, but the final eligibility decision has contractual consequences. Review the table and choose the best ownership for Step 3.

Step	Work performed	Risk / reversibility	Current proposal	What follows
1	Extract member fields	Reversible; checked later	AI	Validation
2	Apply documented formula	Computable; reproducible	AI + code	Review
3	Determine eligibility exception	High stakes; exception can be hard to unwind	AI	Notification
4	Approve exception	Accountable decision	Human reviewer	Notification
5	Notify member	Externally consequential	Human	—

A. AI, because Step 2 is reliable and supplies the needed facts.

B. Collaborative, because AI can prepare the exception analysis but a named person should judge it.

C. AI, but only when the model reports high confidence.

D. Human, with no AI contribution allowed anywhere in the exception analysis.

Answer: B

29. A finance team is testing a redesigned monthly close. The table shows the baseline and two candidate workflows. The business cares primarily about accuracy, not raw speed. Which option should be selected based on the evidence?

Workflow	Time/run	Revision rounds	Manual steps	Accuracy exceptions
Baseline	52 min	2.1	14	4/100
Candidate A	31 min	2.0	8	9/100
Candidate B	38 min	1.2	9	3/100
Candidate C	28 min	3.4	6	11/100

A. Candidate A, because it has the best balance of time and manual effort.

B. Candidate B, because it improves time while reducing revisions and has fewer accuracy exceptions than baseline.

C. Candidate C, because minimizing time is always the first optimization goal.

D. Candidate C, because six manual steps is the strongest evidence of quality.

Answer: B

30. A customer support workflow has been running for six months. The table shows when failures appeared and the cheapest confirmation test. The current incident is wrong only after very long conversations, while a fresh session is correct. Which diagnosis best matches the table?

Observed symptom	Timing	Cheap confirmation	Likely diagnosis
Missing requirements	First response	Compare prompt to required inputs	Under-specification
Format drift	After long session	Repeat in fresh session	Context overload
One arithmetic error	Every run	Run with code execution	Wrong feature
Old product terms	After months	Check source date	Stale configuration

A. Under-specification, because all output failures ultimately originate in prompts.

B. Stale configuration, because any degradation indicates the knowledge source changed.

C. Context overload, because the symptom appears after accumulated conversation and disappears in a fresh session.

D. Wrong feature/model, because a longer context requires a larger model tier.

Answer: C

31. A team is deciding where to put four fixes. The table distinguishes the nature of each correction. Which row is correctly classified?

Fix	Nature	Candidate home	Why
Always state the intended audience in line one	Rule	Standing instruction	Applies every run
Use the current brand voice guide	Reference	Knowledge base	Changes source material
Generate the report in five fixed steps	Procedure	Skill	Encodes ordered steps
Do not send until a named reviewer approves	Gate/control	Knowledge base	Actually a control/gate, not reference content

A. Row 1 only

B. Rows 1 and 2 only

C. Rows 1, 2, and 3 only

D. All four rows

Answer: C

32. A service team compares its old and new workflows. The table records three cycles. The team optimizes for consistency because the output is customer-facing. Which conclusion is best?

Cycle	Old consistency score	New consistency score	Old time	New time
1	82	94	45	31
2	84	93	44	30
3	81	95	46	29
Average	82.3	94.0	45.0	30.0

A. The new workflow is supported by the relevant metric because consistency improved materially across all observed cycles.

B. The new workflow should be rejected because time also decreased, indicating under-review.

C. The old workflow should remain because only one metric can ever be optimized.

D. No conclusion is possible because a baseline is unnecessary for customer-facing work.

Answer: A

33. An operations map was reviewed after volume doubled. The table captures what changed. Which issue most directly explains why a previously valid collaborative gate is now failing?

Control	Before	Now	Observed signal
Review volume	20/week	50/week	Reviewer spends seconds per case
Standing instruction	Current	Current	No drift observed
Map owner	Named	Named	Still available
Knowledge source	Current	Current	No stale-source finding

A. The knowledge source is stale.

B. The map owner is no longer accountable.

C. The gate's staffing/capacity no longer matches the new workflow volume.

D. The standing instruction must be removed because review volume increased.

Answer: C

34. A customer asks whether a workflow is “fully automated.” The table lists the actual controls. Which description is most accurate?

Step	AI role	Human control	External effect
Extract fields	Produces draft	None	None
Flag deviations	Produces candidate flags	Reviewer checks flagged items	None
Draft response	Produces draft	Support lead approves	None
Send response	No AI ownership	Named signer executes	Customer receives notice

A. Fully automated because AI performs the majority of preparation.

B. AI handles the workflow end-to-end, with humans available if something goes wrong.

C. AI assists with extraction, flagging, and drafting, while named human checkpoints control review and external release.

D. Human-only because any human checkpoint means AI contributes no operational value.

Answer: C

35. A team has a workflow incident and reconstructs the following facts. The table shows that the workflow produced 180 reports while using an outdated reference file. Which immediate operating decision is most justified?

Fact	Observed value	Impact	Immediate concern
Reports generated since stale file introduced	180	Potentially affected population	Scope
Reports externally distributed	132	Customer/external exposure	Recipients
Current workflow status	Paused	Further spread stopped	Maintain pause
Known affected source	Product definitions	Content may be wrong	Replace/audit source
Reviewer at release gate	Named, but review did not detect stale source	Gate exists but failed to catch drift	Map/configuration review

- A. Resume immediately because the reviewer exists.
- B. Scope the 132 distributed reports, identify recipients/impact, communicate through the required path, and update the map/configuration before resuming.
- C. Delete the 180 reports without determining which were distributed.
- D. Only correct future reports because prior outputs are outside the workflow.

Answer: B