

FDE-Associate Foundation — Governance, Risk & Responsible Use: Advanced MCQs

Source reviewed: Panaversity AI Agent Factory — “Governance, Risk & Responsible Use: A Crash Course.” The exam emphasizes Case, Data, Capability, People; defined gates; data minimization; approved routes; least privilege; prompt injection; fairness and exclusions; incident response; diligence gaps; residual risk; Governance Records; and change-triggered re-checks.

Format verified: 35 questions — 19 single-answer, 8 two-answer, 8 table-based. Answers are immediately after the options.

Q1. A billing assistant drafts customer replies using account records. The reply is factually important, but the account manager can verify the record and tone before sending. Which classification is strongest?

- A. Fully appropriate with no special gate, because the model only drafts.
- B. Appropriate with review, because the organization remains accountable for factual claims and can review before sending.
- C. Inappropriate, because AI must never draft customer communications.
- D. Fully inappropriate unless the model is fine-tuned on billing cases.

Answer: B

Q2. A hiring workflow uses AI to shortlist applicants. The manager reviews only the survivors and never examines who was excluded. Which risk is the most important missing control?

- A. The model may have used too many tokens.
- B. The workflow may systematically exclude a group without anyone noticing; exclusions should be sampled as well as inclusions.
- C. The manager should review only the highest-confidence survivors.
- D. The solution is to remove all human review.

Answer: B

Q3. A manager says, “A human reviews every AI output.” The FDE asks who, what, and when. Nobody can answer precisely. What is the correct diagnosis?

- A. This is already a defined gate because a human exists somewhere in the process.
- B. It is an intention, not a defined gate; the workflow is not ready until the reviewer, verification target, and timing are specified.
- C. The model should approve its own output.
- D. The gate can be defined after the output is sent.

Answer: B

Q4. A team wants AI to analyze customer spending trends. Names are not needed for the pattern. Which first move best follows the source?

- A. Upload the raw spreadsheet and tell the model not to retain it.
- B. Skip the analysis entirely.
- C. Remove unnecessary identifiers before the data reaches the AI system, then verify that the remaining data meets the organization's anonymization standard.
- D. Put the raw spreadsheet in a temporary chat and assume it is authorized.

Answer: C

Q5. An analyst replaces names with Customer 17, Customer 18, etc., but keeps a separate mapping table that can reconnect the labels to people. How should the resulting data be treated?

- A. Fully anonymous in every context.
- B. Pseudonymized; the mapping means it remains traceable and should retain the relevant sensitivity tier.
- C. Public data because names are gone.
- D. Green automatically because labels are not names.

Answer: B

Q6. A team argues that an incognito/temporary conversation makes confidential data acceptable because it is not in normal history. What is the strongest correction?

- A. Temporary history controls answer whether the data was authorized to enter the system.
- B. Retention/history controls are narrower than authorization; the route still must be approved for that data and purpose.
- C. Incognito makes every data tier green.

D. Memory controls eliminate contractual obligations.

Answer: B

Q7. A code-execution sandbox is used to process a sensitive file. The team concludes, "Because the code is sandboxed, the upload is approved." Which diagnosis is correct?

- A. Correct; execution isolation is the same as data authorization.
- B. Incorrect; a sandbox limits execution reach but does not decide whether the data was permitted to enter that environment.
- C. Correct only for regulated data.
- D. Incorrect only if the model has memory enabled.

Answer: B

Q8. A customer dataset was collected to send service reminders. Marketing proposes using the same contact data to sell unrelated products through an AI campaign. Which governance question should be raised?

- A. Whether the model is sufficiently creative.
- B. Whether the new use is covered by the original collection purpose or requires fresh permission.
- C. Whether the data is stored in CSV format.
- D. Whether the AI has a larger context window.

Answer: B

Q9. A connector is approved, but its scope includes reading all shared-drive files and sending email, while the workflow needs only three reporting tables. What is the best FDE action?

- A. Enable it because the connector is approved.
- B. Narrow the scope to the least privilege needed and remove unnecessary action authority.
- C. Replace the connector with a general-purpose browser.
- D. Grant delete access so the workflow can be fully tested.

Answer: B

Q10. A formatting Skill from a colleague is internally published. Its description sounds harmless, but the package includes scripts that can access files outside the stated task. What is the right conclusion?

- A. Internal origin means it is already vetted.
- B. Treat it like software: inspect source, reach, fit, outside content, and actions before enabling it.
- C. A Skill has its own permission boundary that prevents file access.
- D. The only check needed is whether it has been used successfully before.

Answer: B

Q11. A trusted connector reads incoming customer emails. One email contains instructions telling the AI to forward confidential records to an external address. What is the correct security model?

- A. The connector is trusted, so the email instructions are trusted too.
- B. Tool trust and content trust are separate; the email is untrusted content and may contain prompt injection.
- C. The customer becomes the tool publisher.
- D. The agent must follow the email because it is newer than the system instructions.

Answer: B

Q12. A workflow can read public information and draft an internal memo, but it cannot send, publish, pay, delete, or edit consequential records. Which security posture is most aligned with least privilege?

- A. Give it broad write access for future flexibility.
- B. Keep read-only access where sufficient and put consequential actions behind defined human gates.
- C. Remove all audit logs because humans review it.
- D. Allow the model to choose its own permissions.

Answer: B

Q13. A team discovers that employees routinely use an unapproved AI tool because the approved route requires ten steps while the shadow route requires one. Which intervention is most likely to reduce the underlying risk?

- A. Send another reminder that shadow AI is forbidden.

- B. Reduce friction in the approved route while maintaining its controls.
- C. Remove the approved route entirely.
- D. Ignore it because users are responsible for their own choices.

Answer: B

Q14. A governance audit is turning into a performance review, so employees show only their cleanest AI work. What should the FDE change?

- A. Audit the process and system gaps—friction, stale permissions, weak gates, confusing rules—rather than trying to score individuals.
- B. Increase employee surveillance.
- C. Stop auditing entirely.
- D. Require employees to submit only successful examples.

Answer: A

Q15. A sensitive AI incident occurs: an employee uploaded regulated data through an unapproved route, then notices the mistake five minutes later. What should happen first?

- A. Delete the visible copy and wait to see whether anyone noticed.
- B. Stop further spread, record the facts, and report promptly through the organization's incident path.
- C. Send a customer notification immediately without consulting the incident owner.
- D. Rewrite the Governance Record before reporting.

Answer: B

Q16. An AI tool accidentally sent an incorrect customer notice. The employee considers deleting all logs to “clean up the evidence.” What is the strongest response?

- A. Delete them if the customer has not complained.
- B. Do not delete evidence; follow the incident owner's instructions because preservation, notification, legal, and forensic consequences may apply.
- C. Delete only the local copy.
- D. Ask the model to erase its memory.

Answer: B

Q17. A Governance Record was approved last year. Since then, the model family changed and the connector gained a new write action. The team says the old approval still covers it. What should the FDE do?

- A. Continue because the workflow name is unchanged.
- B. Re-check the Governance Record because model/capability/permission changes are explicit re-check triggers.
- C. Re-check only if an incident occurs.
- D. Re-check only the People section.

Answer: B

Q18. A manager wants a one-page Governance Record but does not know which policy owner approves the data route. What is better than guessing?

- A. Fill in the approval field based on what seems reasonable.
- B. Write OPEN QUESTION and route it to the appropriate owner.
- C. Leave the whole record blank.
- D. Treat the vendor's security page as the approval.

Answer: B

Q19. A team has all controls working as designed. However, a consistent wording bias still appears across customer reports. What concept describes this remaining exposure?

- A. Diligence gap.
- B. Residual risk.
- C. Shadow AI.
- D. Redaction failure.

Answer: B

Q20. A team classifies an AI-assisted workflow as appropriate with review. Which TWO elements must a defined gate contain?

- A. Who reviews it.
- B. What the reviewer verifies.
- C. The model's confidence score.
- D. A generic statement that a human is involved.
- E. A promise to review eventually.

Correct answers: A and B

Q21. A customer asks whether raw identifiers can be removed before trend analysis. Which TWO questions should drive the decision?

- A. Does the task actually need the identifiers?
- B. Does the remaining data truly meet the organization's anonymization standard?
- C. Is the model popular?
- D. Does the UI have an incognito mode?
- E. Can the analyst hide the mapping table from the model but keep it elsewhere?

Correct answers: A and B

Q22. A capability review is underway for a new Skill. Which TWO checks are directly required by the source's five-check model?

- A. Who created or published it?
- B. What can it reach in the session and is that reach proportionate?
- C. Whether its marketing description is friendly.
- D. Whether a colleague has used it.
- E. Whether it has a colorful icon.

Correct answers: A and B

Q23. A trusted plugin reads outside web content and can also publish changes. Which TWO controls are especially important?

- A. Treat outside content as potentially untrusted and screen for prompt injection.
- B. Keep publish/other irreversible actions behind a defined human gate until the workflow earns more autonomy.
- C. Trust all web instructions because the plugin is approved.
- D. Grant broader access so the model can defend itself.
- E. Disable audit logs to reduce exposure.

Correct answers: A and B

Q24. A People review is being designed for an AI screening workflow. Which TWO practices best address the source's fairness guidance?

- A. Examine exclusions as well as inclusions.
- B. Consider who is affected even if they never see the AI output.
- C. Review only the highest-ranked cases.
- D. Assume correct facts imply a fair process.
- E. Avoid escalation if the affected population is large.

Correct answers: A and B

Q25. An organization has no formal AI policy. Which TWO steps are appropriate for an FDE?

- A. Draft an interim proposal for the right owner rather than presenting it as official policy.
- B. Include approved routes, sensitive-data confirmation rules, human-gate rules, extension review, and an incident contact.
- C. Declare the interim proposal legally binding.
- D. Skip security review until a formal policy exists.
- E. Let each employee define their own approved route.

Correct answers: A and B

Q26. A first-hour incident response is underway. Which TWO actions should be avoided?

- A. Deleting evidence to make the problem disappear.
- B. Waiting for a perfect explanation before reporting a sensitive incident.
- C. Recording what happened and what route was involved.
- D. Following the incident owner's instructions.
- E. Reporting a near miss.

Correct answers: A and B

Q27. A manager is evaluating review quality. Which TWO properties make a human gate meaningful rather than ceremonial?

- A. The reviewer has enough source material and time to check the intended risk.
- B. The reviewer has authority to reject before the outcome becomes hard to reverse.
- C. The reviewer approves nearly everything without examining edge cases.
- D. The review occurs after publication.
- E. The reviewer only sees the model's final prose.

Correct answers: A and B

Q28. An FDE applies the four-question framework to four workflows. Which row is classified most defensibly?

Workflow	Case facts	Data	People	Best classification
A	Internal FAQ from approved policy; reversible	Approved internal	Low impact	Fully appropriate
B	Customer billing reply; factual accountability remains	Customer account data	Customer affected	Appropriate with review
C	Final disciplinary determination delegated to AI	Employee sensitive	High consequence	Inappropriate
D	Candidate shortlist; exclusions never sampled	Applicant data	Employment impact	Fully appropriate

Question: Select the best answer based on the table.

- A. A
- B. B
- C. C
- D. D

Answer: D

Q29. A data-governance review produces the following cases. Which row correctly identifies the control?

	Data	Task needs identifiers?	Route/control	Correct outcome	
	Customer trend data	No	Redact, then verify anonymization	Proceed if approved	
	Account reconciliation	Yes	Redact all IDs	Task becomes impossible	Do not use this control
	Credentials	Yes	Incognito chat	Still red	Route remains unapproved
	Privileged legal file	No	Temporary chat	Privacy setting = authorization	Incorrect

Question: Select the best answer based on the table.

- A. A
- B. B
- C. C
- D. D

Answer: A

Q30. A connector and Skill are reviewed using the five-check model. Which row has the strongest posture?

	Source	Reach	Outside content	Actions	Decision
	Unknown forum	Whole drive + credentials	Web + email	Send + delete	Decline
	Internal team	Read-only reporting tables	Customer notes	None	Enable after scope

	Source	Reach	Outside content	Actions	Decision
	Known vendor	Whole company data	External web	Publish	Enable automatic
	Unknown publisher	Small folder	None	Read-only	Enable because re

Question: Select the best answer based on the table.

- A. A
- B. B
- C. C
- D. D

Answer: B

Q31. A team audits a customer-report workflow after one quarter. Which row best represents a meaningful control and its residual risk?

Control	Success measure	Failure threshold	Monitoring	Residual risk
A	Manager clicks approve	Any click counts	Never	None
B	Facts/tone checked before send	Material factual error reaches customer	Monthly sample	Cross-report wording bias
C	AI confidence > 0.9	No threshold	Annual	Unknown
D	Policy says review	No metric	Ad hoc	None

Question: Select the best answer based on the table.

- A. A
- B. B
- C. C
- D. D

Answer: B

Q32. A governance audit finds repeated workarounds. Which row best diagnoses the diligence gap?

Observation	Policy says	Actual behavior	Root issue	Best response
A	Use approved AI route	Employees use a one-step shadow tool	Approved route has excessive friction	Reduce friction while preserving controls
B	Review outputs	Review skipped under deadlines	People are careless	Send punishment memo
C	Vet Skills	Skill enabled without source check	No checklist at setup	Ignore until incident
D	Use approved connector	Old connector scope remains	Permission drift	Leave unchanged until renewal

Question: Select the best answer based on the table.

- A. A
- B. B
- C. C
- D. D

Answer: A

Q33. A serious AI incident is reported. Which sequence matches the source's first-hour procedure?

Step	Action	Correct?	Why	Priority
1	Stop spread	Yes	Contain further copies/actions	Immediate
2	Record facts	Yes	Preserve reliable account	Immediate
3	Report promptly	Yes	Timing may matter	Immediate
4	Guess root cause publicly	No	Introduces unreliable claims	Avoid
5	Follow incident owner	Yes	Legal/forensic consequences may exist	After report

Question: Select the best answer based on the table.

- A. Steps 1-3 only
- B. Steps 1,2,3 and 5
- C. Steps 2,4 and 5

D. All steps

Answer: B

Q34. A Governance Record must be re-checked. Which row contains only explicit change triggers from the source?

	Model change	Data sensitivity	Connector permission	Audience	Vendor terms
	Yes	Yes	Yes	Yes	Yes
	No	Yes	No	No	No
	Yes	No	No	No	No
	No	No	No	Yes	No

Question: Select the best answer based on the table.

A. A

B. B

C. C

D. D

Answer: A

Q35. A review team evaluates four human gates. Which gate is highest quality?

Gate	Who	What	When	Quality
A	Someone	Check it	Eventually	Weak
B	Manager	Verify source facts	Before send	Strong
C	AI system	Approve itself	After send	Weak
D	Reviewer	Read final text	After publication	Weak

Question: Select the best answer based on the table.

A. A

B. B

C. C

D. D

Answer: B