

FDE Associate Foundation — Code You Never Write

Advanced scenario-based MCQ examination • Source: Panaversity “Code You Never Write: A Crash Course”

Verified structure: 35 questions = 19 single-answer + 8 two-answer + 8 table-based. Questions emphasize FDE reasoning, commissioning, verification, troubleshooting, surface selection, and safe execution.

1. A finance team asks an AI, “Look at our 9,000-row expense export and tell me roughly how spending changed by month.” The answer looks plausible, but no code block or row count is shown. The monthly figures will drive a budget revision. What is the strongest next move?

- A. Accept the figures because the wording says “roughly.”
- B. Ask the AI to summarize the largest month only.
- C. Explicitly require code to run, show the code, and provide row count/date-range checks before analysis.
- D. Ask the AI to explain why its estimates are probably accurate.

Answer: C

2. A clinic manager wants a monthly no-show report. The task is repeated, contains thousands of appointment rows, and a missed row could distort staffing decisions. Which classification is most defensible?

- A. Answer problem, because the AI can reason about patterns.
- B. Code problem, because Repetition and Precision independently trigger computation.
- C. Answer problem unless the file has more than 10,000 rows.
- D. Code problem only if the manager knows Python.

Answer: B

3. An FDE is writing a brief for a school export. The AI knows how to process CSV files, but it does not know that “Pending” means “not yet submitted” in this school. Where should that domain fact primarily go?

- A. Goal
- B. Input
- C. Rules
- D. Output

Answer: C

4. A customer says, “Find students who did not submit.” The export contains duplicate submissions, blank names, and submissions from people not on the class list. Which brief design most reduces silent damage?

- A. Ask the AI to use its best judgment on unusual rows.
- B. Specify duplicate semantics and state exactly how unmatched and unreadable rows must be handled.
- C. Tell the AI to ignore all malformed rows without reporting them.
- D. Require the AI to infer missing policy from the student names.

Answer: B

5. A customer asks for an exact payroll total. The AI returns a clean number and a narrative but no code block. What should the FDE infer?

- A. Code probably ran invisibly, so the number is sufficient.
- B. The AI may have estimated or skimmed; absence of the requested proof means computation has not been established.
- C. The number is trustworthy if it has two decimal places.
- D. The correct response is to read the generated narrative more carefully.

Answer: B

6. A script processes 1,400 transactions but the report contains 1,381 rows. The customer knows all 1,400 should be represented. What is the best diagnostic prompt?

- A. "Fix the script immediately."
- B. "Rewrite the whole solution in another language."
- C. "Show the missing 19 rows and delete any suspicious records."
- D. "Report the row counts, date range, and representative input/output rows; explain where the 19 disappeared before changing logic."

Answer: D

7. A reconciliation script has produced the expected total for a known closed month. The customer now wants to trust it on a full year. Which verification step most directly establishes that the logic generalizes from a known case?

- A. Ask the AI to make the code shorter.
- B. Run a known-answer test on another independently known slice.
- C. Ask the AI whether it feels confident.
- D. Replace the script with manual arithmetic for the full year.

Answer: B

8. A high-stakes report will be sent to an external auditor. The script passes a known-answer test, reality checks, plain-English replay, and adversarial review. Which additional check is most aligned with the course's highest verification tier?

- A. Ask the same AI to repeat the answer three times.
- B. Run the same brief and data through a second model family and compare results.
- C. Hide the input data and see whether the AI can reconstruct it.
- D. Remove the known-answer sample so the model cannot overfit.

Answer: B

9. A script fails with a red error message after a file export changes. The FDE does not understand the technical error. What is the intended operating pattern?

- A. Diagnose the error manually before involving the AI.
- B. Paste the complete error, ask the AI to diagnose, fix, and rerun.
- C. Delete the script and start over immediately.
- D. Ignore the error if the previous month's report worked.

Answer: B

10. A report completes without an error, but sales are \$4.2M when the customer expects roughly \$12M. Which response best distinguishes a wrong-but-running result from a simple crash?

- A. State the symptom and expected value, then inspect processed row count, date range, and sample rows before fixing.
- B. Tell the AI "the answer is wrong" and demand a rewrite.
- C. Increase the model temperature so it considers more possibilities.
- D. Ask for a prettier chart to expose the problem.

Answer: A

11. The same repair has failed three times and each attempt produces another variant of the same wrong result. What does the source recommend?

- A. Continue making smaller edits until the existing approach works.
- B. Stop, restate the problem fresh, and request two fundamentally different approaches.
- C. Delete the known-answer sample.
- D. Disable all validation checks so the script can finish.

Answer: B

12. A customer has a task they will perform every Monday. The first AI run works. What turns that one-off solution into a durable operational asset?

- A. Keep only the final report.
- B. Save the script, brief, and known-answer sample together with a clear description of inputs and rules.
- C. Ask the AI to remember the task without files.
- D. Recreate the prompt from memory every Monday.

Answer: B

13. A team is choosing a surface for a one-time analysis of three CSV files on a borrowed laptop. They do not need to touch local files. Which is the strongest choice?

- A. Browser chat sandbox
- B. Terminal agent with broad filesystem access
- C. Desktop agent configured to reorganize the laptop
- D. Scheduled local worker

Answer: A

14. A customer's workflow reads files already stored in a local folder, runs every morning, and must touch network devices. Which surface fits the operational need best?

- A. Browser sandbox
- B. Claude Code/OpenCode-style agent in the local environment
- C. Static HTML report
- D. One-off chat with manually uploaded files

Answer: B

15. A knowledge worker wants an agent to reorganize a folder but is uncomfortable with direct file mutation. Which capability most directly addresses that concern?

- A. Larger context window
- B. Plan-then-approve before actions execute
- C. Automatic retries without review
- D. More programming languages

Answer: B

16. An FDE is about to run an unproven cleanup script on a customer's entire Documents directory. Which change most directly reduces blast radius before execution?

- A. Run it faster so the window for failure is smaller.
- B. Point it at the smallest relevant folder, work on copies, require a dry run, and write outputs to new files.
- C. Ask the agent to be careful.
- D. Increase filesystem permissions so the agent can self-repair.

Answer: B

17. A customer wants code to "compute which employees should be promoted." The data is clean and the arithmetic is trivial, but promotion criteria involve judgment and fairness. What is the most important boundary?

- A. Code should make the final promotion decisions because arithmetic is objective.
- B. Code may compute relevant measures, but the human decision criteria and judgment remain outside one-prompt computation.
- C. The task is safe if the model explains its reasoning.
- D. A second model should automatically make the final decision.

Answer: B

18. A customer asks an agent to submit a tax filing automatically every quarter. The agent can prepare the numbers reliably, but the submission is effectively irreversible once sent. What is the best architecture?

- A. Fully unattended submission because repetition makes it a code problem.
- B. Let code prepare the filing, then require a human to review and submit it.
- C. Avoid code entirely and calculate everything manually.
- D. Let a second AI submit after checking the first AI.

Answer: B

19. An FDE notices a customer repeatedly uploads 40 files every week to browser chat and then manually downloads the script. The workflow is correct and verified. What is the strongest graduation signal?

- A. The model has become less intelligent.
- B. Uploading has become the annoying part of the week; move the workflow to a local folder-based surface.
- C. The customer should keep uploading because browser sandboxes are always safer.
- D. Replace computation with manual spreadsheets.

Answer: B

20. A customer wants to trust an AI-generated reconciliation that will affect an external audit. Which TWO actions are most defensible according to the source?

- A. Run a known-answer test on an independently known slice.
- B. Ask the AI to repeat the same result until it is consistent.
- C. Perform a cross-model check using the same brief and data.
- D. Remove edge-case rules to simplify the computation.
- E. Skip human review if all totals match.

Correct answers: A and C

21. A brief for messy data is failing silently. Which TWO additions most directly reduce ambiguity rather than merely adding technical detail?

- A. Explicit Rules containing domain-specific facts.
- B. A preferred Python library name.
- C. Explicit Edge cases describing how imperfect rows are handled.
- D. A request for more colorful charts.
- E. A requirement to use recursion.

Correct answers: A and C

22. A running result is suspicious. Which TWO pieces of evidence are especially useful before changing the code?

- A. Processed row count and date range.
- B. The AI's confidence score alone.
- C. Representative rows as the code sees them.
- D. A longer prompt demanding perfection.
- E. A second visualization with the same underlying data.

Correct answers: A and C

23. A script will rename and delete files on a local machine. Which TWO controls most directly constrain the blast radius before action?

- A. Work on a backup/copy first.
- B. Give the agent the whole home directory so it can see everything.
- C. Require a dry-run list and approval before mutation.

■ D. Overwrite originals so there is one canonical copy.

■ E. Disable logs to reduce clutter.

Correct answers: A and C

24. A team wants to preserve a reusable workflow six months from now. Which TWO artifacts should be kept together with the script?

■ A. The plain-English brief describing what and why.

■ B. Only the latest output report.

■ C. A small sample whose correct answer is known.

■ D. A screenshot of the chat interface.

■ E. A model-generated confidence score.

Correct answers: A and C

25. A private customer export contains names and account numbers that are irrelevant to the computation. Which TWO actions best align with the source?

■ A. Create a cleaned copy with unnecessary identifiers removed.

■ B. Upload the raw file first and redact later.

■ C. Work only on the cleaned copy and confirm what was removed.

■ D. Assume a browser sandbox makes sensitive data automatically safe.

■ E. Disable all verification to avoid exposing data in prompts.

Correct answers: A and C

26. A customer has tried the same repair three times and the approach still fails. Which TWO actions best follow the recommended troubleshooting pattern?

■ A. Treat three failures as evidence the approach may be wrong.

■ B. Keep changing one character at a time indefinitely.

■ C. Restate the problem fresh and request two different approaches.

■ D. Remove the known-answer test.

■ E. Increase the file size to stress the same approach.

Correct answers: A and C

27. A customer wants to use AI-written code for a high-consequence workflow. Which TWO statements capture the course's trust model?

■ A. Confidence should come from the size and quality of the checks, not from how convincing the answer sounds.

■ B. Seeing a code block proves the result is correct.

■ C. Code is exact enough to be exactly wrong, so independent verification remains necessary.

■ D. If two models agree, omitted business rules can no longer be a problem.

■ E. A successful run eliminates the need for human review in high-stakes work.

Correct answers: A and C

28. A payroll FDE is evaluating four proposed workflows. The customer wants the FDE to classify each by whether one-prompt code is inside the course's safe territory. Use the table, not general intuition.

Workflow	Signals / characteristics	Consequence	Human judgment required	Proposed classification
A	2,000-row monthly payroll reconciliation; exact totals; repeatable	High precision	Review claims before payroll release	Code problem with verification
B	Draft a persuasive memo explaining a policy	Primarily language/judgment	Substantive editorial judgment	Answer problem
C	Compute promotion decisions from employee scores using a disputed fairness rule	Precise arithmetic but contested criteria	Yes; criteria determine outcome	Code alone is insufficient
D	Rename 500 archived files using a deterministic naming rule	Volume + Files	Low after dry run	Code problem; safe with blast-radius controls

- A. Use the first row as the most defensible decision.
- B. Use the second row as the most defensible decision.
- C. Use the third row as the most defensible decision.
- D. Use the fourth row as the most defensible decision.

Answer: D

29. An FDE must choose a verification strategy for four outputs. The team has limited time but cannot treat precision-critical results casually. Which row represents the strongest minimum verification match?

Work item	Stakes	Known answer available?	Required verification shown	Best match
A	Personal spending curiosity	No	Reality questions	Appropriate
B	Colleague-facing report	Yes	Known-answer test + plain-English replay	Appropriate
C	External audit reconciliation	Yes	All five checks + human review	Appropriate
D	Health-related calculation	Yes	Only visual inspection of chart	Appropriate

- A. Use the first row as the most defensible decision.
- B. Use the second row as the most defensible decision.
- C. Use the third row as the most defensible decision.
- D. Use the fourth row as the most defensible decision.

Answer: C

30. A customer reports that an analysis is wrong, but the script completed successfully. The FDE wants to diagnose rather than guess. Which row shows the best evidence path?

Case	Observed output	First diagnostic evidence	Likely hidden failure to test	Best next move
A	\$4.2M vs expected ~\$12M	Processed rows, date range, first five rows	Partial sheet/export	Inspect before rewriting
B	Red traceback	Full error text	Column/file incompatibility	Paste error; diagnose/fix/rerun
C	Same fix fails 3 times	Three failed iterations	Wrong approach	Fresh restatement + two alternatives
D	Correct closed-month total	Exact match	None indicated	Delete the known-answer sample

- A. Use the first row as the most defensible decision.
- B. Use the second row as the most defensible decision.
- C. Use the third row as the most defensible decision.
- D. Use the fourth row as the most defensible decision.

Answer: A

31. A team is deciding which surface to use. The same conceptual task can be executed on different surfaces, but reach and persistence differ. Which row is the most defensible recommendation?

Need	Environment	Repetition	Special reach	Recommended surface
A	One-off CSV analysis	Borrowed laptop	None	Browser sandbox
B	Daily network-device health report	Local machine	Schedule + network access	Folder-based local agent
C	Folder reorganization with approval before mutations	Desktop workstation	Human approval	Plan-then-approve desktop app
D	One-off file analysis	Local machine	None	Scheduled worker

- A. Use the first row as the most defensible decision.
- B. Use the second row as the most defensible decision.
- C. Use the third row as the most defensible decision.
- D. Use the fourth row as the most defensible decision.

Answer: B

32. A photo-cleanup workflow can delete files permanently. The customer asks why the FDE insists on a staged process. Which row best represents the intended safety sequence?

Step	Action	Mutation?	Review gate	Purpose
A	1	Backup originals	No	Recover from bad run
B	2	Generate full deletion plan	No	See exact changes before execution
C	3	Approve reviewed deletion set	Yes	Human confirms mutation
D	4	Write results over originals	Yes	Preserve canonical state

- A. Use the first row as the most defensible decision.
- B. Use the second row as the most defensible decision.
- C. Use the third row as the most defensible decision.
- D. Use the fourth row as the most defensible decision.

Answer: D

33. A reusable script is being handed to another analyst. The goal is to prevent silent rule drift and preserve the ability to re-verify. Which row contains the strongest artifact set?

Package	Contents	Known sample	Domain rules preserved	Re-run readiness
A	Script only	No	Partly	Low
B	Brief + script	No	Yes	Medium
C	Brief + script + known-answer sample	Yes	Yes	High
D	Latest HTML report + screenshot	No	No	Low

- A. Use the first row as the most defensible decision.
- B. Use the second row as the most defensible decision.
- C. Use the third row as the most defensible decision.
- D. Use the fourth row as the most defensible decision.

Answer: C

34. An FDE is evaluating whether a project belongs inside the “code you never write” territory. Which row is the clearest boundary case that should be escalated beyond one-prompt code?

Project	Core operation	Consequence / autonomy	Boundary assessment	Why
A	Merge exports and chart monthly spend	Compute/report	Reversible	Inside
B	Find duplicate photos and propose deletions	Compute + propose	Deletion requires approval	Inside with safety gate
C	Automatically email customers every week without oversight	Unattended external action	Runs while nobody watches	Beyond one-prompt territory
D	Calculate a teacher-specific grade	Compute under explicit rules	Human owns grading policy	Inside if rules are explicit

- A. Use the first row as the most defensible decision.
- B. Use the second row as the most defensible decision.
- C. Use the third row as the most defensible decision.
- D. Use the fourth row as the most defensible decision.

Answer: C

35. A security-conscious FDE is reviewing four proposed local execution plans. Which row most completely follows the four blast-radius rules?

Plan	Backup/copy	Dry run + approval	Smallest folder	New output files
A	Yes	Yes	Yes	Yes
B	Yes	No	Yes	Yes
C	No	Yes	Yes	Yes
D	Yes	Yes	No	Yes

- A. Use the first row as the most defensible decision.
- B. Use the second row as the most defensible decision.
- C. Use the third row as the most defensible decision.
- D. Use the fourth row as the most defensible decision.

Answer: A