

FDE Associate Foundation — Code You Never Write

35 Scenario-Based MCQs | 19 single-answer • 8 two-answer • 8 table-based

Based primarily on the complete Code You Never Write crash course: VPRF classification, explicit computation, five-section briefs, verification ladder, failure diagnosis, reusable scripts, execution surfaces, blast-radius controls, project selection, and the boundary between code, products, automation, and human judgment.

1. A finance team asks an AI to answer, “What did we spend on travel this year?” The workbook contains 18,000 transaction rows, and the resulting figure will be used in a board pack. The model can summarize the workbook directly, but the team wants the cheapest reliable approach. What should the FDE recommend?

- A. Treat it as a code problem because Precision and Files fire, explicitly require computation, and verify the result against a known answer.
- B. Treat it as an answer problem because the question is phrased conversationally and the model can summarize spreadsheets.
- C. Ask for an approximate total first and only use code if the approximation differs materially from expectations.
- D. Tell the model to use Python immediately, because the FDE should choose the implementation language for financial tasks.

Answer: A

2. A customer wants to know whether a single paragraph is persuasive. There is no file processing, repeated computation, or numerical precision requirement. The customer insists that the AI “write and run code” because they believe code is always more reliable. What is the best FDE response?

- A. Agree, because running code is safer than reasoning for every task.
- B. Explain that this is an answer problem and that forcing code would add unnecessary machinery rather than improve the judgment.
- C. Convert the paragraph into a CSV so that the Files signal fires.
- D. Build a reusable script first because every customer problem should become an asset.

Answer: B

3. An operations analyst uploads a spreadsheet and asks, “What does this data show?” The model returns several exact-looking totals but provides no code block, row count, column names, or date range. The totals look plausible. What is the most defensible conclusion?

- A. The result is trustworthy because plausible totals are sufficient for exploratory analysis.
- B. The model probably used code internally, so the missing code block is only a presentation issue.
- C. Treat the response as a possible glance-based estimate and explicitly force computation plus evidence that the file was processed.
- D. Replace the model immediately because any missing code block proves the model cannot work with files.

Answer: C

4. A team is commissioning a monthly reconciliation. The developer-like user writes: “Use pandas, loop through every row, normalize the strings, then join on transaction ID.” The business owner actually cares about matching rules, exclusions, output, and what to do with malformed rows. What is the best FDE correction?

- A. Add more implementation details so the model cannot improvise.
- B. Replace all technical instructions with a one-line goal and let the model infer everything else.
- C. Keep the technical plan but add a second prompt describing the business outcome.
- D. Brief the problem using Goal, Input, Output, Rules, and Edge cases, leaving implementation choices to the AI.

Answer: D

5. A school administrator asks an AI to find students who have not submitted an assignment. The class list has 180 students; the submission file has duplicate submissions and some names contain extra spaces. Which brief detail most directly prevents a false “submitted” result caused by duplicate rows?

- A. State that a student who submitted twice still counts as submitted once.
- B. Ask the AI to sort the file alphabetically before processing.
- C. Tell the AI to use Python because duplicate handling requires code.
- D. Ask for an HTML output because HTML makes duplicates visible.

Answer: A

6. A payroll workflow produces a precise total for 9,000 employees. The AI shows a successful run and a polished report, but the FDE has no independently known payroll subtotal to compare against. What is the strongest next step before using the number?

- A. Ask the AI to rate its confidence from 1–10.
- B. Create or select a small known-answer test and require the computed result to match that independently known answer.
- C. Run the same prompt again in the same conversation and accept matching results.
- D. Inspect the formatting of the final report because formatting consistency is evidence of computational correctness.

Answer: B

7. A script completes successfully but reports annual sales of \$4.2M. The finance lead knows the year should be near \$12M. The FDE does not want to guess at the cause or immediately rewrite the code. What should they ask for first?

- A. “Fix the code and rerun it using a different library.”
- B. “Start over from scratch because the number is wrong.”
- C. “Show the processed row count, date range, and first five rows as the code sees them, then investigate the discrepancy.”
- D. “Run the same code on a second model and choose whichever number is larger.”

Answer: C

8. An AI-generated reconciliation script fails three times after the team makes small modifications to the same approach. Each run produces a different error, and the conversation has become long and confused. What does the course's troubleshooting rule imply?

- A. Continue iterating until the original approach finally works.
- B. Add more technical implementation details to force convergence.
- C. Delete the error messages so the model can reason without distraction.
- D. Treat three failed attempts as evidence the approach is wrong; restate the problem fresh and request two different approaches.

Answer: D

9. A teacher has built a script that checks assignment submissions every month. The script, brief, and a small known-answer sample are stored together. Six months later, the teacher receives a new class list. What is the main value of having kept the script?

- A. The teacher can rerun the solved workflow instead of commissioning the entire process from scratch.
- B. The script guarantees that future input files cannot contain malformed data.
- C. Keeping the script means verification is no longer necessary on new data.
- D. The script converts the workflow into a judgment problem rather than a code problem.

Answer: A

10. A professional has 40 files to analyze once and is using a browser AI chat with a temporary sandbox. They are considering installing a local coding agent solely for this one-off task. What does the course suggest?

- A. Always install a local agent when more than ten files are involved.
- B. Stay on the browser surface for one-off exploratory work; local agents become valuable when files live on the machine or the task repeats.
- C. Use a desktop agent because temporary sandboxes cannot run Python.
- D. Avoid code entirely because uploading many files automatically makes the task unsafe.

Answer: B

11. A network engineer needs a job to run every morning against devices on the company's private network without manually uploading files. Which execution surface best fits the requirement?

- A. Browser chat sandbox, because it is the safest surface regardless of connectivity.
- B. Desktop app with plan-then-approve, because scheduling is handled by the approval step.
- C. An agent running on the engineer's machine, such as a terminal-based coding agent, because it can reach local/network resources and be scheduled.
- D. A reusable browser Skill, because Skills eliminate the need for local execution.

Answer: C

12. A knowledge worker wants an AI to reorganize a folder containing sensitive documents. The proposed action will rename, move, and potentially delete files. Which control sequence best follows the course's blast-radius guidance?

- A. Let the agent operate on the original folder because backups slow down automation.
- B. Give the agent broad permissions first, then reduce them after observing one successful run.
- C. Ask the agent to explain its code syntax before it makes changes.
- D. Work on the smallest folder, make a backup/copy first, require a dry-run plan and approval, and write outputs to new files where possible.

Answer: D

13. A customer asks an AI to “compute which employees should be promoted” from performance data. The arithmetic is straightforward, but promotion depends on criteria involving leadership potential, context, and organizational judgment. Which boundary from the course applies?

- A. Code should compute the data-driven measures, while the human retains the judgment decision.
- B. Code should make the promotion decision because judgment becomes more reliable when expressed numerically.
- C. The task is always an answer problem because promotion is a human resource decision.
- D. A second model should make the final decision so that two independent judgments cancel bias.

Answer: A

14. A team is deciding whether to build a small custom script or buy an existing application for a workflow. The workflow follows the same standard rules used by nearly every organization in the industry, and a mature application already handles it. What is the strongest recommendation under the course's “off-the-shelf-app test”?

- A. Build the script anyway because custom code is more flexible.
- B. Prefer the existing application; the course's project territory is strongest for market-of-one problems whose rules are specific to the user.
- C. Build the script only if the application has a graphical interface.
- D. Use AI code because any task involving files automatically qualifies as a custom project.

Answer: B

15. A sensitive HR export contains names, account numbers, and salary figures, but the analysis only needs department and salary bands. The organization's policy permits AI analysis only on minimized data. What should the FDE do before uploading?

- A. Upload the full export because the model needs context to avoid errors.
- B. Encrypt the file inside the prompt and assume the model can safely decrypt it.
- C. Create a cleaned copy containing only necessary fields, confirm what was removed, and use an approved tool/data policy.
- D. Convert the file to HTML because HTML is safer than CSV.

Answer: C

16. A code run returns a wrong number with no error. The team says, "The AI is broken." Which response best matches the course's diagnostic posture?

- A. Ask the AI to apologize and regenerate the report.
- B. Immediately switch programming languages.
- C. Repeat the same request three times to see whether the average is stable.
- D. Describe the discrepancy as a symptom, state the expected value, and ask the AI to inspect row counts, date range, and representative rows before changing anything.

Answer: D

17. A company wants to automatically submit tax filings every quarter using AI-written code. The calculations are highly accurate in tests, but the submission has no practical undo once transmitted. Which architecture is most aligned with the course's stated boundary?

- A. Let the code prepare the filing, but require a human to review and perform the final submission.
- B. Fully automate submission because accuracy testing eliminates the need for human involvement.
- C. Use a cross-model check and then allow automatic submission if both models agree.
- D. Move the task into a browser sandbox because temporary environments are safer.

Answer: A

18. A product manager wants to build a permanent web dashboard that hundreds of customers log into, with accounts, payments, uptime requirements, and support obligations. The initial prototype can be produced with one AI prompt. What should the FDE recognize?

- A. The dashboard remains a simple script because AI wrote the first version.
- B. This has crossed the boundary into a real product requiring software engineering discipline, rather than remaining a one-prompt code task.
- C. The correct response is to keep regenerating the dashboard in chat until it is stable.
- D. A known-answer test alone is sufficient for production because the dashboard is generated code.

Answer: B

19. A high-stakes external reconciliation will be sent to an auditor. The first AI's output passes a known-answer test, reality checks, plain-English replay, and an adversarial pass. The team asks whether any further validation is useful. What does the course recommend for this level of stakes?

- A. No further validation is ever useful once the first model passes its checks.
- B. Repeat the same model in a fresh conversation and accept agreement.
- C. Run the same brief and data through a second AI family and compare results, while remembering both can share the same missing-rule mistake.
- D. Read every line of generated code manually before any further work.

Answer: C

20. A team is triaging requests before deciding whether to use AI-written code. Which TWO conditions are enough to classify a task as a code problem under the course's VPRF rule?

- A. The task contains professional judgment.
- B. The task involves enough items that doing it by hand is impractical.
- C. A wrong digit could have meaningful consequences.
- D. The user prefers Python.
- E. The task is described in a long prompt.
- F. The task is conversational rather than file-based.

Correct answers: B and C

21. A customer is writing a five-section code brief for a messy spreadsheet workflow. Which TWO sections are specifically intended to capture information that a generic technical implementation would not safely infer?

- A. Goal
- B. Input
- C. Rules
- D. Output
- E. Edge cases
- F. Programming language

Correct answers: C and E

22. A precision-critical report is about to be delivered. Which TWO practices provide the strongest evidence that the computation actually processed the intended data?

- A. Ask for the exact row count, column names, and date range before analysis.
- B. Accept the result because the final chart looks polished.
- C. Run a known-answer test on a subset whose correct result is independently known.
- D. Ask the model to promise that it used code.
- E. Tell the model to use Python regardless of the output type.
- F. Repeat the same calculation in the same context until the number feels stable.

Correct answers: A and C

23. A script is wrong but produces no error. Which TWO diagnostic actions are directly aligned with the course's recommended response?

- A. State the symptom and the value you expected.
- B. Delete the current script and start coding from scratch immediately.
- C. Ask for processed row count, date range, and representative rows as the code sees them.
- D. Focus first on whether the AI used Python or JavaScript.
- E. Ignore the discrepancy if the number is directionally plausible.
- F. Ask the model to increase its confidence score.

Correct answers: A and C

24. A team is moving a recurring workflow from browser chat to a local agent. Which TWO conditions most strongly justify the move?

- A. Uploading the same files has become an ongoing annoyance.
- B. The team wants the AI to choose a prettier chart.
- C. The task needs to access files already on the machine or local network.
- D. The task has no repetition and uses only one tiny file.

■ E. The browser sandbox is more secure than all local execution, so migration is required.

■ F. The team wants to avoid verification.

Correct answers: A and C

25. A local agent is about to modify customer files. Which TWO safeguards most directly reduce blast radius before the first real run?

■ A. Give it the entire company drive so it can resolve missing references.

■ B. Make a backup/copy and operate on the smallest necessary folder.

■ C. Require a dry-run plan and human approval before destructive changes.

■ D. Disable all output files so the agent has nowhere to write.

■ E. Ask it to explain the programming syntax instead of the planned changes.

■ F. Run it unattended once so the agent can learn the safe behavior.

Correct answers: B and C

26. A team wants confidence in a high-stakes result produced by AI-written code. Which TWO statements correctly reflect the course's verification ladder?

■ A. Cross-model agreement is useful additional evidence, but both models can share the same missing-rule mistake.

■ B. A known-answer test is unnecessary if the code ran without errors.

■ C. Plain-English replay can expose wrong business assumptions even when the evaluator cannot read code.

■ D. A polished HTML report is equivalent to a verification test.

■ E. Mutation testing means adding more production data until the output stabilizes.

■ F. An adversarial pass replaces the need for any known-answer test.

Correct answers: A and C

27. A customer asks an AI to perform a no-undo action with serious external consequences. Which TWO design choices fit the course's stated boundary?

■ A. Let code prepare the action while a human performs the final irreversible step.

■ B. Automate it if two models produce identical outputs.

■ C. Treat judgment calls as computation whenever the inputs are numerical.

■ D. Keep human approval where the act cannot practically be undone.

■ E. Move the task to a temporary sandbox and then execute it automatically.

■ F. Skip dry runs because external systems provide their own audit logs.

Correct answers: A and D

28. An FDE is classifying five customer requests before choosing the appropriate workflow. Use the evidence in the table and select the request whose classification is wrong.

Request	Volume/Precision	Files/Repetition	Current classification
Draft a polite decline email	Low	No	Answer
Total 2,000 sales rows by month	High	Yes	Code
Judge whether a business idea is good	Low	No	Answer
Rename 300 photos from embedded dates	High	Yes	Code
Summarize one news article	Low	No	Code

A. Draft a polite decline email

B. Total 2,000 sales rows by month

C. Rename 300 photos from embedded dates

D. Summarize one news article

Answer: D

29. A payroll FDE is reviewing two candidate briefs. The business requires a precise result and has messy input data. Based on the table, which brief is stronger?

Section	Brief A	Brief B	Assessment
Goal	Calculate payroll	Identify exact payroll total for approved period	B is more precise
Input	Payroll file	CSV export, ~9,000 rows, date range stated	B specifies input
Output	Number	Number plus reconciliation summary	B specifies deliverable
Rules	Use normal payroll logic	State overtime, refunds, exclusions, and period rules	B captures business rules
Edge cases	Handle errors	Skip/list unreadable rows and flag malformed amounts	B handles imperfect data

A. Brief A, because it is shorter and leaves the AI more freedom.

B. Brief B, because it specifies business rules and imperfect-data behavior without dictating implementation.

C. Brief A, because technical detail is always safer for precision tasks.

D. Neither, because code briefs should not include edge cases.

Answer: B

30. A team is validating an AI-generated sales report before delivery. The table shows the verification evidence gathered. Which claim is most defensible?

Check	Result	Implication	Review status
Known-answer test	Passed March subtotal	Basic computation agrees with independent knowledge	Pass
Rows in vs. rows out	9,000 in / 9,000 represented	No obvious row-loss signal	Pass
Plain-English replay	Explains exclusions and refund handling	Business logic can be inspected without code syntax	Pass
Adversarial pass	Found one assumption about date boundaries	Potential risk surfaced for review	Needs confirmation
Cross-model check	Second model agrees	Independent confirmation, but shared-rule errors remain possible	Supporting evidence

A. The result is guaranteed correct because all checks passed.

B. The report is unusable because the adversarial pass found an assumption.

C. The evidence is strong, but the surfaced date-boundary assumption still needs a human judgment/confirmation before high-stakes use.

D. Only the cross-model result matters because two models are more reliable than one.

Answer: C

31. A team has an incident in which a script runs successfully but returns a suspicious total. The table records the first diagnostic evidence. Which next action is best?

Evidence	Observed	Likely direction	Next focus
Error message	None	Wrong-but-running	Inspect inputs/output
Expected annual total	~\$12M	Output is ~\$4.2M	Find discrepancy
Processed rows	3,100	Source has ~9,000	Possible truncation
Date range detected	Jan–Apr	Source is Jan–Dec	Possible incomplete input
Sample rows	Correct format	No obvious parsing error	Continue scope investigation

- A. Rewrite the algorithm immediately because the total is wrong.
- B. Investigate why only 3,100 rows and Jan–Apr were processed before changing the computation logic.
- C. Ignore the row count because successful execution proves the file was fully read.
- D. Run the same script three more times and average the totals.

Answer: B

32. A department is deciding which execution surface to adopt. The table captures the workflow requirements. Which surface is the best fit?

Requirement	Browser sandbox	Local coding agent	Desktop plan-then-approve
One-off uploaded file	Strong	Possible	Possible
Files already in local folder	Weak	Strong	Strong
Scheduled daily run	Weak	Strong	Strong/depends on setup
Human approval before file changes	Possible	Possible	Strong
Zero installation	Strong	Weak	Weak

- A. Browser sandbox, because it is strongest on zero installation.
- B. Local coding agent, because every workflow should run on the user's machine.
- C. Desktop plan-then-approve, because all work involves file changes.
- D. Choose based on the workflow: browser for one-off work, local agent for persistent/scheduled folder work, and desktop approval surfaces when plan-before-action is the priority.

Answer: D

33. A local cleanup job proposes deleting duplicate files. The FDE reviews the safety evidence below. Which decision is most appropriate?

Control	Status	Finding	Decision
Backup copy	Complete	Originals preserved	Satisfied
Dry-run list	Complete	Shows every proposed deletion	Satisfied
Human review	Pending	One near-duplicate may be a different photo	Must review
Scope	Small folder	Limited blast radius	Good
Output	New cleaned copy	Originals untouched	Good

- A. Approve deletion immediately because the backup exists.
- B. Reject the workflow permanently because duplicate detection is imperfect.
- C. Review the near-duplicate and approve only after the human gate is satisfied.
- D. Remove the backup because it is redundant with the dry run.

Answer: C

34. A company is deciding whether to automate an external action. The table describes four candidates. Which candidate should remain human-controlled under the course's boundary?

Action	Computation difficulty	Undo available?	Judgment/stakes
Monthly internal chart	Low	Yes	Low
Rename local copies	Medium	Yes	Low
Prepare tax filing	High	No	High
Calculate monthly totals	High	Yes	Low

- A. Monthly internal chart
- B. Rename local copies

- C. Prepare tax filing
- D. Calculate monthly totals

Answer: C

35. A team has a solved recurring workflow and is deciding what to preserve for future runs. The table shows the artifacts available. Which package best follows the course's “describe once, run forever” principle?

Artifact	Present?	Purpose	Keep?
Original brief	Yes	Preserves Goal/Input/Output/Rules/Edge cases	Yes
Reusable script	Yes	Executes the solved workflow	Yes
Known-answer sample	Yes	Regression/verification reference	Yes
Temporary chat only	Yes	Contains latest conversation	No
Human memory	Yes	Informal recollection of assumptions	No

- A. Keep the brief, reusable script, and known-answer sample together; do not rely on chat history or memory alone.
- B. Keep only the temporary chat because it contains everything.
- C. Keep only the script because the code contains all business rules.
- D. Keep only the known-answer sample because the model can infer the workflow from it.

Answer: A