

FDE-Associate Foundation Exam Preparation

Just Delegate It: A Crash Course — 30 Application-Focused MCQs

Source: Agent Factory / Panaversity — “Just Delegate It: A Crash Course”

<https://agentfactory.panaversity.org/docs/just-delegate-it-crash-course>

Questions are original scenario-based applications of concepts from the source; source wording is not reproduced.

Section A — Scenario-Based MCQs (1–20)

Question 1

A product team asks an AI assistant, “What are some good AI-agent courses?”

The response is a list, but the team still has to compare providers, check whether the courses are current, and decide what to hand to the learner. The team wants the assistant to perform the first-pass work and return something they can inspect.

Which change most directly converts this interaction from asking into delegation?

- A. Ask the assistant to make its answer longer and include more background.
- B. Ask the assistant to produce a defined result, such as a comparison table with links and stated checking requirements.
- C. Ask the assistant to explain why each course might be useful before listing them.
- D. Ask the assistant to provide its internal reasoning before giving the course list.

Correct Answer: B

Why: Delegation is distinguished by asking for a result that can be checked and used, rather than leaving the work as advice.

Question 2

An FDE is designing a workflow for an AI assistant to compare several software vendors. The FDE knows the desired output but does not want to dictate every search step, page visit, or comparison action because the assistant may find a better route. Which instruction best reflects the course’s approach to delegation?

- A. Specify every click and page in the exact order the assistant must follow.
- B. Describe the finished result and important limits, while leaving routine method choices to the assistant.
- C. Tell the assistant to imitate the FDE’s previous workflow exactly.
- D. Avoid constraints so the assistant has unlimited freedom.

Correct Answer: B

Why: The course emphasizes describing the outcome and leaving routine method selection to AI, while preserving meaningful constraints.

Question 3

A junior engineer writes a request for an AI research task but leaves unclear what success looks like, what information the assistant already needs, what it may decide, and how the output will be checked. The engineer wants a reusable work order rather than a vague prompt. Which set of fields should be completed before the assistant begins?

- A. Model, temperature, token limit, browser, API key, and system prompt.
- B. Outcome, context, constraints, authority, deliverable, and verification.
- C. Question, keywords, search engine, citations, summary, and opinion.
- D. Input, code, deployment, logging, monitoring, and rollback.

Correct Answer: B

Why: Those six fields constitute the Delegation Brief described in the course.

Question 4

A company needs a report containing the current prices of three AI tools.

An assistant gives plausible prices but provides no sources and says the facts are based on what it “knows.” The prices will be used for a purchasing decision next week, so stale information would materially affect the result. What should the engineer do first?

- A. Accept the prices because the assistant sounds confident.
- B. Ask for current research and verify a source and its date for at least the claims that matter.
- C. Ask the assistant to repeat the same answer with more confidence.
- D. Remove the price requirement and use approximate historical values.

Correct Answer: B

Why: Current facts require looking now, and the course stresses checking the source and date rather than trusting a citation or confident wording.

Question 5

An FDE has a private internal policy document that an AI assistant cannot find on the public web. The FDE needs answers based only on that document and does not want the assistant to fill gaps from memory. Before sharing the material, the FDE also needs to reduce exposure of private information. Which approach best fits the course?

- A. Paste the entire document and ask the model to use both the document and its general knowledge.
- B. Remove unnecessary private details, provide the relevant source, instruct the assistant to stay inside it, and check quoted support.
- C. Summarize the policy from memory and ask the model to improve the summary.
- D. Upload the document without any source-boundary instruction because the model will infer the rule.

Correct Answer: B

Why: The source should be provided, the job should be constrained to that source, and private information should be removed before pasting.

Question 6

A support manager asks an assistant to “summarize this incident.” The manager actually needs a reusable one-page incident report with sections for impact, timeline, unresolved items, and source caveats. A paragraph in chat would force the manager to restructure it manually. Which design choice best demonstrates the deliverable principle?

- A. Ask only for a concise answer and format it later.
- B. Name the exact usable artifact and its required structure, then check the output against that structure.
- C. Ask the assistant to decide whether a deliverable is necessary after responding.
- D. Avoid specifying the format so the assistant can be creative.

Correct Answer: B

Why: A deliverable is a usable thing—document, table, file, etc.—and its form should be specified and checked.

Question 7

An AI agent is preparing a public website for an organization. The agent may draft the page, but publishing it would create a public commitment that the owner must approve. The owner wants the assistant to work independently on routine drafting while retaining control over the consequential action. Which authority boundary is most appropriate?

- A. Let the agent publish immediately because it can undo the change later.
- B. Allow drafting and preparation, but require explicit approval before publication.
- C. Prevent the agent from doing any drafting at all.
- D. Allow publication whenever the agent believes the page is high quality.

Correct Answer: B

Why: The course uses boundaries such as “draft, but do not send” and “build, but wait before publishing” to preserve human control over consequential actions.

Question 8

A recurring weekly report requires the same background documents and the same standing rule: any unsupported claim must be marked “unverified.” The analyst has started pasting the same material into every chat and wants to reduce repetition without losing the standing context. What is the most appropriate architectural move described in the course?

- A. Put the repeated source files into a Project as knowledge and the persistent rule into standing instructions.
- B. Put everything into the model's hidden memory and stop verifying outputs.
- C. Rewrite the entire prompt every week with more examples.
- D. Use a new temporary chat each week so no context can persist.

Correct Answer: A

Why: Projects hold persistent knowledge and standing instructions for repeated work; the course distinguishes these from today's job in the chat.

Question 9

An assistant returns a vendor comparison, but one column uses the wrong language and one course is outside the user's stated constraints. The assistant has not yet produced a completely unusable result. The engineer wants the smallest correction that gets the job back on course instead of rewriting the whole brief. Which intervention is most appropriate?

- A. Restart the entire project from scratch every time.
- B. Name the wrong direction and redirect the assistant toward the required constraints.
- C. Accept the result and manually repair every row.
- D. Immediately reject all AI assistance for future work.

Correct Answer: B

Why: When the work is heading the wrong way but is not yet wrong, the course recommends redirecting rather than restating the entire job.

Question 10

A model repeatedly struggles with a multi-step calculation even after the engineer has clearly stated the desired outcome and constraints. The job is important, and the assistant's current capability appears insufficient for the required level of reasoning. Which response best matches the course's escalation strategy?

- A. Add many unrelated instructions until the model succeeds.
- B. Escalate to a stronger thinking/model setting if available, or break the job into smaller checked steps.
- C. Remove verification so the assistant can finish faster.
- D. Keep sending the identical message indefinitely.

Correct Answer: B

Why: Escalation can mean a stronger model/thinking setting; if unavailable, make the job smaller and check each step.

Question 11

A finance team receives an AI-generated calculation that will appear in an audit-critical report. The assistant cites a source, but the team has not checked the underlying number, the source date, or whether the source actually supports the claim. Which verification sequence best follows the course?

- A. Read the answer twice, then accept it if it sounds internally consistent.
- B. Identify the material claim, trace it to the supporting source and date, challenge how it could be wrong, then decide what action to take.
- C. Check only whether a citation symbol appears beside the claim.
- D. Ask the assistant whether it is confident and use its confidence as the final test.

Correct Answer: B

Why: Verification is Identify, Trace, Challenge, Decide; a citation alone is not verification.

Question 12

A project lead has verified that an AI-generated recommendation contains one specific factual error that can be fixed without changing the entire job. The source for the correct value is available, and the rest of the deliverable is sound. What should the lead do?

- A. Correct the specific error and verify the corrected result.
- B. Reject every part of the deliverable automatically.
- C. Delete the source because it caused confusion.
- D. Ask the assistant to produce a completely unrelated recommendation.

Correct Answer: A

Why: A specific fixable error maps to the Correct decision, followed by checking the actual corrected result.

Question 13

An AI assistant has produced a recommendation based on a secondary blog. The recommendation affects a regulated process, and the organization's official documentation is available. The blog and official documentation disagree. Which verification approach is most consistent with the course?

- A. Prefer whichever source is newer without examining the content.
- B. Trace the claim to the primary source, check the date, and assess whether the evidence supports the recommendation.
- C. Average the two sources and use the midpoint.
- D. Accept the blog because secondary sources are easier to understand.

Correct Answer: B

Why: The course says a provider's own/primary page generally beats a blog about the provider, while also requiring attention to dates and possible change.

Question 14

A team has corrected the same AI result three times. Each revision improves the output less than the previous one, and the task concerns a specialized field where mistakes have meaningful consequences. The team is considering another prompt rewrite. What does the course suggest this pattern may indicate?

- A. The prompt should always become longer.
- B. The job may need to be taken back by a qualified person rather than repeatedly prompting the model.
- C. Verification should be skipped because the model has seen the feedback.
- D. The result should be published immediately to avoid further iteration.

Correct Answer: B

Why: Diminishing improvement after repeated corrections is a signal to raise the reviewer to a person with domain expertise.

Question 15

An organization wants to run the same delegation brief through a second AI assistant. The team's purpose is to test whether the brief travels and to observe differences in the runs, not to create a permanent ranking of models. Which procedure best matches the course?

- A. Rewrite the brief for the second assistant so it matches that product's controls.
- B. Paste the complete brief unchanged, record what happened, and compare observations without declaring a winner.
- C. Use only the second assistant's result because it is newer.
- D. Change the brief until both assistants produce identical wording.

Correct Answer: B

Why: A good brief describes the job rather than the tool; the course recommends recording observations rather than ranking assistants.

Question 16

A scheduled AI job produces a weekly market summary while nobody is watching. The first run looks polished and contains links, but no one has checked whether the important claims are still accurate. The team wants to use the automation responsibly. What is the key lesson to apply to the first unattended result?

- A. A scheduled job is automatically trustworthy because it ran successfully.
- B. Apply the same verification discipline to the result; a job running is not the same as a job working.
- C. Remove all source checking because scheduled tasks are designed for autonomy.
- D. Let the assistant decide whether its own output is safe to publish.

Correct Answer: B

Why: The course explicitly warns that an automated run still requires verification; execution does not equal correctness.

Question 17

An engineer asks an assistant to research a rapidly changing product feature. The assistant provides a link, but the page has no obvious date and the cited sentence does not actually support the assistant's claim. The engineer wants to know whether the research requirement was genuinely satisfied. Which conclusion is best supported?

- A. The presence of a link proves that web research occurred correctly.
- B. The citation is insufficient; the engineer must inspect the supporting sentence and relevant date before accepting the claim.
- C. The assistant should be trusted because the claim is plausible.
- D. Any undated page should automatically be treated as false.

Correct Answer: B

Why: The course distinguishes a citation from verification and requires checking the supporting sentence and date when currency matters.

Question 18

An FDE is deciding whether an AI-connected system should be allowed to act on an external service. The agent needs access to company data and may create or modify resources. Before enabling the connection, the FDE wants to reason about access rather than simply turning on the connector. Which set of questions best matches the course?

- A. Is the interface attractive, is the model popular, and is the response fast?
- B. What may it read, what may it do, can the information source be trusted, and is the organization allowed to grant that access?
- C. How many tokens does the connector consume and what color is its icon?
- D. Can the agent make the decision without any human involvement?

Correct Answer: B

Why: The course identifies four pre-action access questions: read permissions, action permissions, source trust, and whether the access is allowed.

Question 19

A team is preparing a public legal communication with AI assistance. The draft looks clean and has passed several automated checks. The team wants to know whether a person still needs to review it before publication. Which course rule directly applies?

- A. No person is needed if the model has a high-quality setting.
- B. Public or legal communications with lasting consequences require a named person to read the result before release.
- C. Only the first draft ever needs human review.
- D. Human review is required only when the assistant reports uncertainty.

Correct Answer: B

Why: The course lists public/legal communications among outputs that should not leave without a person reviewing them.

Question 20

A delegation brief produces a good result, but the engineer notices that a failure occurred because a constraint was omitted. The engineer fixes the result manually and moves on. The same failure is likely to recur next week. Which process improvement best follows the course?

- A. Keep the fix private so the next run is independent.
- B. Add the learned constraint or authority boundary to the brief so the same drift is prevented next time.
- C. Remove the entire task from the assistant.
- D. Increase the answer length without changing the brief.

Correct Answer: B

Why: The course says intervention should be followed by writing the lesson into the brief as a constraint or authority line.

Section B — Two-Selection MCQs (21–25)

Question 21

An FDE is building a Delegation Brief for a recurring research task. The team wants the brief to be checkable and to make human responsibility clear. Which TWO elements below belong directly in the six-line brief described by the course? Select exactly TWO options.

- A. Authority describing what AI may decide without asking.
- B. Verification describing what will be checked before acceptance.
- C. A mandatory list of every browser click the AI must perform.
- D. A permanent ranking of the AI models the company uses.

Correct Answers: A and B

Question 22

A current-facts research task is being delegated for a purchasing decision. The engineer wants to distinguish genuine research from a confident answer based on model memory. Which TWO practices are most directly supported by the course? Select exactly TWO options.

- A. Require current information and inspect an accessible source supporting an important claim.
- B. Check the relevant page date when currency matters.
- C. Treat any citation marker as proof that the claim is verified.
- D. Prefer an answer with no sources if it is concise.

Correct Answers: A and B

Question 23

An AI assistant is allowed to prepare a deliverable but must not make a consequential external commitment without approval. The FDE is designing the authority boundary before enabling a connector. Which TWO controls best preserve the intended human decision point? Select exactly TWO options.

- A. Allow the assistant to draft or prepare the action.
- B. Require explicit approval before the consequential action occurs.
- C. Allow the assistant to silently change the boundary when it finds a faster route.
- D. Remove all verification because approval itself is sufficient.

Correct Answers: A and B

Question 24

A recurring task is moving from a one-off chat to a reusable operating pattern. The same background sources and standing rules appear each time, and the team wants AI to receive them automatically. Which TWO actions align with the course? Select exactly TWO options.

- A. Store persistent source material as Project knowledge.
- B. Store persistent rules as standing instructions.
- C. Delete the source after each run to force a fresh start.
- D. Put every changing detail into standing instructions.

Correct Answers: A and B

Question 25

An AI result has been reviewed and one claim is materially important but poorly supported. The reviewer cannot yet determine whether the claim is true. The reviewer wants to follow the course's verification method rather than guess. Which TWO actions are appropriate? Select exactly TWO options.

- A. Trace the claim to the source and examine the supporting evidence.
- B. Investigate further when the evidence remains incomplete or contradictory.
- C. Accept the claim because it has a citation.
- D. Reject the entire workflow immediately without checking the claim.

Correct Answers: A and B

Section C — Detailed Table-Based MCQs (26–30)

Question 26

An engineering manager is diagnosing where a delegation workflow failed. The table below shows what happened at different stages. For each row, consider the symptom and the course principle that should guide the next action. Which overall response is the most appropriate solution strategy?

Observed situation	What the team saw	Needed response
Start	The request names a topic but not a usable result.	Clarify the outcome and deliverable.
During work	The assistant follows a poor route but has not produced anything yet.	Use a different method.
After work	A material claim has a source but the source is stale.	Take and challenge the claim before acceptance.
Next run	The same background and rules are pasted again.	Move persistent context into a Project.

- A. Treat every row as a prompt-writing problem and rewrite the entire job each time.
- B. Match the intervention to the symptom: define the result, redirect drift, verify material claims, and persist repeated context.
- C. Accept the result whenever the assistant provides a source, then automate the next run.
- D. Put every possible instruction into Project knowledge and eliminate human review.

Correct Answer: B

Question 27

A company is deciding how much autonomy to grant an AI agent that can access external systems. The table describes four proposed boundaries. Based on the course's authority principles, which implementation strategy should the team adopt?

Proposed action	Boundary	Reason in the workflow
Research	May research sources.	No external commitment is created.
Purchase	May buy items automatically.	Money movement may be difficult to reverse.
Draft	May prepare a message.	The draft can be reviewed before release.
Send	May send externally without approval.	Sending creates a visible commitment.

- A. Allow all four actions because the agent is supervised at the system level.
- B. Permit routine research and drafting, but require approval before money movement or external sending.
- C. Block research and drafting but allow purchases if the source is reputable.
- D. Allow sending automatically and review the message only after it is delivered.

Correct Answer: B

Question 28

An analyst receives an AI-generated report and must decide where to spend verification effort. The table lists candidate claims and risk characteristics. Using the course's review-threshold approach, which verification plan is most appropriate?

Claim	Stakes	Reversibility	Audience/Exposure
-------	--------	---------------	-------------------

A	Low	Easy to undo	Internal analyst
B	High	Hard to undo	Client
C	Medium	Easy to undo	Internal team
D	High	Hard to undo	Public/regulatory context

- A. Verify only A because it is easiest to inspect.
- B. Prioritize B and D because high stakes and low reversibility/audience or regulatory exposure push them onto the must-check list.
- C. Verify C only because medium-risk claims are the most common.
- D. Verify all claims equally and spend the same effort regardless of consequences.

Correct Answer: B

Question 29

An FDE compares two runs of the same brief. No wording in the brief changed, but the world and the assistant's behavior may have changed. The table shows the differences. Which interpretation and next step best follows the course?

Difference	Possible explanation	What to do
Course closed	World changed	Check the current source/date.
Different source selected	AI behavior changed	Record the observation and assess the claim.
Different output format	AI behavior changed	Compare against the deliverable requirement.
Different price	World or source timing changed	Trace the current price before accepting it.

- A. Declare one assistant the winner because its result looks cleaner.
- B. Record the observations, distinguish world changes from AI differences, and verify material claims using the same method.
- C. Rewrite the brief until both assistants produce identical results.
- D. Ignore the second run because model differences make verification unnecessary.

Correct Answer: B

Question 30

A supervisor is reviewing a proposed final deliverable before release. The table summarizes four types of outputs. The supervisor must decide which workflow should be treated as requiring named human review every time, even if the AI run looked clean. Which option identifies the correct general strategy?

Output	Consequence if wrong	Human review requirement
Internal brainstorming	Limited	Routine review
Final client deliverable	Material	Named person reviews before release
Audit-critical calculation	Material/financial	Named person reviews before release
Public legal communication	Potentially lasting	Named person reviews before release

- A. Require human review only for brainstorming because it is the least structured.
- B. Require named human review for final client deliverables and other high-consequence outputs listed by the course.
- C. Remove human review whenever the assistant provides citations.
- D. Review only outputs that contain an explicit "unverified" label.

Correct Answer: B